An Evaluation on the Performance of Code Generated with WebAssembly Compilers

Raymond Phelan

A dissertation submitted in partial fulfilment of the requirements of
Technical University Dublin for the degree of
M.Sc. in Computing (TU060)

Date: June 14, 2021

Declaration

I certify that this dissertation which I now submit for examination for the award of MSc in Computing (Stream), is entirely my own work and has not been taken from the work of others save and to the extent that such work has been cited and acknowledged within the text of my work.

This dissertation was prepared according to the regulations for postgraduate study of the Technical University Dublin and has not been submitted in whole or part for an award in any other Institute or University.

The work reported on in this dissertation conforms to the principles and requirements of the Institute's guidelines for ethics in research.

Signed:

Raymond Phelan

Date: June 14, 2021

Abstract

WebAssembly is a new technology that is revolutionizing the web. Essentially it is a low-level binary instruction set that can be run on browsers, servers or stand-alone environments. Many programming languages either currently have, or are working on, compilers that will compile the language into WebAssembly. This means that applications written in languages like C++ or Rust can now be run on the web, directly in a browser or other environment. However, as we will highlight in this research, the quality of code generated by the different WebAssembly compilers varies and causes performance issues.

This research paper aims to evaluate the code generated by a number of existing WebAssembly compilers in order to determine whether or not there is a significant difference in their performances regarding execution times.

Keywords: WebAssembly, WebAssembly Compilers, Benchmarking Performance, WASM, WAT, AssemblyScript, C/C++, Rust

Acknowledgments

I would like to thank my supervisor, Paul Kelly, for his on going support and advise throughout this dissertation. I also express my sincere gratitude to all the lecturers in TU Dublin that I have had the pleasure of attending their classrooms.

I would also like to say a very special thank you to, my mother June Phelan, my father Raymond Patrick Phelan, my wife Diana Phelan and my children, Lucas Perico Phelan, Tawana Akinlolu and Tamiya June Phelan, for their constant encouragement and support.

Finally, I would like to thank my classmates that I have had the pleasure of getting to know along this journey.

Contents

List of Figures

List of Tables

Listings

List of Acronyms

WASI	WebAssembly System Interface
WAT	WebAssembly Text Format
AS	AssemblyScript
VM	Virtual Machine
EVM	Ethereum Virtual Machine
CSV	Comma Separated Values
CPU	Computer Processing Unit

Chapter 1

Introduction

1.1 Background

WebAssembly, also known as WASM, is a new technology revolutionizing the web. It is a low-level binary instruction set, similar to pure assembly language, designed to run at near-native speeds otherwise only achieved with C/C++ code. It can be run on the client-side in web-browsers, on the server-side with NodeJS, and even as a stand-alone environment beyond the web, known as WASI (WebAssembly System Interface) [1].

WebAssembly was first announced on December 5th, 2019, by the World Wide Web Consortium (W3C). For many years, JavaScript has been the only option for providing interactive applications in websites (Musch et al., 2019), although there have been many attempts to remedy this, for example Adobe Flash and Microsoft Active-X. However, these needed to be installed as browser extensions which prevented them from large scale acceptance. WebAssembly on the other hand, requires no extensions to be installed and is already supported by all of the major web-browsers.

Programs written in statically typed languages, such as C, C++, and Rust can now be compiled into WebAssembly. As JavaScript is a dynamically typed language, it cannot be directly compiled into WebAssembly. TypeScript however, a subset of the JavaScript language, can be compiled into WebAssembly. Many other programming

---

[1]https://wasi.dev/

languages are currently working towards enabling compilation to WebAssembly and are currently listed here [2]. This opens the door for many new possibilities as developers from these languages can now develop for the web, where previously this was not possible.

As WebAssembly is essentially binary code, this cannot be easily understood by human readers. There is however a textual representation of WebAssembly, known as WAT (WebAssembly Text Format)[3]. There are tools available for converting pure WebAssembly binary code into WAT format and vice-versa, such as the WABT tool[4]. C and C++ can be compiled to WebAssembly using the Emscripten[5] compiler and Rust can be compiled using the wasm-pack[6] compiler. AssemblyScript[7] is a variant of TypeScript which can compile directly to WebAssembly. These compilers will be the focus of this research project.

1.2   Research Project/problem

As more and more WebAssembly compilers emerge, we can expect to see a large variety of applications developed in multiple programming languages and compiled into WebAssembly. However, these compilers do not necessarily produce the same binary code for a given program written in different programming languages, as we will demonstrate later in this paper. In fact, even for a very simple program, such as a function to calculate the Fibonacci sequence, the different compilers are producing very different WebAssembly binary format code.

These differences can be visually verified by inspecting the WAT format of the WebAssembly function. Since the WebAssembly compilers are producing vastly different code between them, even if they are performing the same function, this may lead to questions such as - which compiler generated the correct code? Do the compiled We-

---

[2]https://github.com/appcypher/ awesome-wasm-langs

[3]https://developer.mozilla.org/en-US/docs/WebAssembly/Understanding$_the_text_format$

[4]https://github.com/WebAssembly/wabt

[5]https://emscripten.org/

[6]https://rustwasm.github.io/wasm-pack/installer/

[7]https://www.assemblyscript.org/

bAssembly modules all have the same performances, or do they vary in efficiency in terms of execution times and file sizes?

Initially, one might have thought that once compiled into WebAssembly, there would be little or no difference in the overall performance between the WebAssembly modules compiled with difference compilers. However, if they have significantly different performances, even though they are all running as WebAssembly files, this means that it is important to understand which programming language and respective compiler developers should choose before developing WebAssembly applications.

These concerns will be the purpose and focus of this research paper. We propose the following research question:

> *Is there a significant difference in the performance of WebAssembly modules, in terms of execution times, for the same program written in different programming languages and compiled into WebAssembly using their respective compilers?*

1.3 Research Objectives

The main goal of this research is to determine whether or not there is a significant difference in the performance of WebAssembly modules generated from different compilers. As we will discuss later in the literature review, at the time of writing this paper, no other research was found that answered our proposed research question. In order to answer the question, we identified the following research objectives:

- Set up a VM (Virtual Machine) environment for running our experiments

- Identify programming languages that can be compiled into WebAssembly

- Set up the respective WebAssembly compilers for each of the selected programming languages

- Identify and code relevant functions and algorithms that can be written in the selected programming languages and that can be properly compiled into We-

3

bAssembly, maintaining the same syntax as much as possible between the languages, and considering the current limitations of WebAssembly and compilers

- Create a benchmark runner for executing the WebAssembly modules from each of the selected programming languages and gather their execution times for post-analysis

- Perform statistical analysis on the gathered execution times in order to answer the research question, by either accepting or rejecting the null or alternate hypotheses.

1.4   Research Methodologies

The type of research we will perform is considered a primary research, since the data that will be used in our experiments does not exist currently in the reviewed literature. The data will be generated by executing the WebAssembly modules in a single environment and gathering their execution times.

In terms of the research objective and whether it implies quantitative or qualitative research, we will be performing our statistical analysis based on numerical data representing the execution times of the WebAssembly modules. Therefore our research can be considered to be a quantitative research.

The form of research we are conducting can be considered to be an empirical research, since we will be making direct observations on the data collected through our scientific experiment. We will also be defining the null and alternative hypotheses based on our research question.

Finally, the reasoning involved in our research can be considered to be deductive reasoning, since we will be answering the research question by testing the hypotheses through a process of evaluating the statistical data obtained during our experiments. We will then draw conclusions to our research question based on the outcome of testing the null and alternative hypotheses.

1.5 Scope and Limitations

The aim of this research paper is to compare the performance of WebAssembly modules generated by different compilers for different programming languages. Our choice of programming languages was influenced by the current available support for each language. We ultimately chose C, C++, Rust and AssemblyScript for our experiments as they provided enough documentation and support for implementing their respective compilers. Long et al. (2021) also came to the same conclusion regarding WebAssembly compilers with sufficient support for different programming languages. There are of course other WebAssembly compilers currently available for compiling other programming languages into WebAssembly, however, they will remain out of scope for this research, but will be suggested as part of future work in this area of research in the concluding chapter.

We also found limited support and resources for implementing benchmarking functions in the languages that we selected, considering that they all needed to maintain the same syntax, and take into account the current limitations of WebAssembly and the compilers. Therefore chose benchmarking functions that could be easily written in all of the selected programming languages.

Another important aspect of this research paper worth highlighting is that we did not necessarily try to find the most efficient implementation for each benchmarking function. Once we were able to replicate the same function in all the selected programming languages, then we would be able to properly compare their performance, regardless of whether the functions were the most efficient implementation or not. Furthermore, this research is limited to execution of WebAssembly modules in a NodeJS environment, rather than the full spectrum of possible environments, such as web-browsers and IoT devices, which we will also suggest as part of future work in the concluding chapter.

1.6   Document Outline

This research project will be divided into logical chapters. Chapter 1 provides some background knowledge on WebAssembly in general. Here we explain the research problem we have identified and present the research question. We also describe the research objectives and methodologies used. Finally we provide an insight into the scope and limitations of this research.

Chapter 2 will involve a review of the current literature based on relevant scientific publications in order to highlight the research gaps and importance of our research question. Chapter 3 describes the experiment design and set up used to gather the required data for our evalutation. Chapter 4 will provide an in-dept evaluation and discussion on the results gathered through our statistical tests. Finally, chapter 5 will provide a conclusion to this research, describing the impact of our contributions and outlining future work.

Chapter 2

Review of existing literature

Although WebAssembly is still relatively new, it is rapidly becoming a highly researched topic in the last 18 months. Researchers have been evaluating the possible benefits and use cases of WebAssembly, while developers have been enjoying freedom of choice in the programming languages they use to develop their applications, using a WebAssembly compiler to achieve the near native speeds promised by WebAssembly.

In this section, we will review currently available literature that carried out some benchmarking on the performance of WebAssembly, or that proposed WebAssembly as solution to improve some existing applications, compared with the traditional approaches. In doing so, we will reveal the gap in the literature that has led us to our research question, highlighting that not enough research and evaluation has been done on the performance of code generated by the various WebAssembly compilers. We highlight the research gap that choosing only one WebAssembly compiler to evaluate the performance of WebAssembly is insufficient.

Sandhu et al. (2018) compared the performance of JavaScript and WebAssembly against C for sparse matrix-vector multiplication. They discover a slowdown of 2.2x to 5.8x with JavaScript versus C, but WebAssembly performing similar or better at times when compared with C in the browser. They used the Emscripten compiler to generate the WebAssembly modules from C code for their experiments. However, only the one WebAssembly compiler was used, therefore it is unknown whether WebAssembly code generated from other compilers might have produced different results.

Sandhu et al. (2020) compared WebAssembly to native C code, again using sparse matrix computations. However, they developed the WebAssembly implementations by hand rather then using a compiler. This time they discover that WebAssembly executed in the Chrome browser has performance issues due to memory addressing in the x86 instruction set. Ultimately, this leaves a gap in the research, as perhaps it would be important to understand if various WebAssembly compilers might have produced more efficient WebAssembly code than the hand-written implementation.

Jangda et al. (2019) use the PolybenchC Benchmark Suite[1] to compare WebAssembly compiled Unix applications versus native speeds inside the browser. Their findings revealed that WebAssembly ran up to 55% slower. However, the WebAssembly modules that were used in their experiments were generated from the Emscripten compiler only, so once again we do not know if WebAssembly from other compilers might result in different performances.

Haas et al. (2017) also compared the performance of WebAssembly using the PolybenchC Benchmark Suite on browser JavaScript engines, such as Mozilla SpiderMonkey, Google V8 and Microsoft Chakra. The generation of WebAssembly modules was done with the assistance of the OCaml Programming Language[2], rather than using any compiler. Once again, there was no evaluation of the WebAssembly code in comparison to what other compilers would generate by implementing different programming languages.

Herrera et al. (2018) evaluate the performance of JavaScript and WebAssembly compared to native C code. They used five of the benchmarks identified in the Ostrich Benchmark Suite (Khan et al., 2015) for numerical computing, and compiled them into WebAssembly modules using Emscripten. Their comparisons were done between web browsers, IoT devices and NodeJS. Their findings once again show that WebAssembly comes close to native C performance, but is faster than JavaScript. However, these experiments also were limited by only evaluating the WebAssembly generated from one compiler, rather than multiple compilers.

---

[1]http://web.cs.ucla.edu/ pouchet/software/polybench/

[2]https://ocaml.org/manual/coreexamples.html

Reiser and Bläser (2017) presented their own WebAssembly compiler, Speedy.js, which compiles JavaScript/TypeScript to WebAssembly, although their compiler uses the Emscripten runtime library. They use a number of computing algorithms, such as Fibonacci, Prime Number and Merge Sort, to evaluate the performance of TypeScript compared to the WebAssembly implementation generated by their proposed solution. Indeed they find WebAssembly performance to be faster than plain TypeScript, however only the WebAssembly from their compiler was evaluated, rather than including WebAssembly generated by other compilers.

Protzenko et al. (2019) also present their own compiler, which compiles Low*, a low-level subset of F* programming language, into WebAssembly. They develop an alternative to the JavaScript encryption library, LibSignal[3], which is used in many modern services such as WhatsApp, Skype and Signal for managing end-to-end encryption. They report no speed up between WebAssembly and the JavaScript library, mainly due to overhead between encoding and decoding between WebAssembly and JavaScript. They suggest that future JavaScript libraries should be designed to be more WebAssembly friendly. However, only WebAssembly modules generated from their own compiler were used in this evaluation, therefore it is unknown if other WebAssembly compilers might have produced more efficient WebAssembly code.

Watt et al. (2019) propose a similar solution to cryptography using WebAssembly, rather than JavaScript. However, they use a different approach, proposing CT-WASM as an extension to WebAssembly, or as a new programming language that is similar to TypeScript, which includes its own compiler for generating pure WebAssembly modules. However, the performance of their generated WebAssembly modules is not evaluated against WebAssembly modules generated from other compilers.

Mendki (2020) proposed a WebAssembly Serverless Edge Computing as an alternative to the native container-based application. They developed a compute and memory intensive application with file I/O and image classification using Rust and compiled it into WebAssembly. They found that WebAssembly had a lower footprint than the container-based solution while also showing faster start-up times, but the runtimes

---

[3]https://github.com/signalapp/libsignal-protocol-javascript

were slower than the native container application. However, only WebAssembly compiled from Rust was evaluated in their proposal, therefore it is unknown how well WebAssembly compiled from other languages might have performed.

A WebAssembly solution to the stateless containers used by existing serverless platforms was published in (Shillaker & Pietzuch, 2020), providing isolated memory and CPU (Computer Processing Unit), with state sharing capabilities. Their solution used the LLVM compiler[4] to translate applications into WebAssembly from multiple programming languages such as, C, C++, Python, JavaScript and TypeScript. Their evaluation reported a 2x performance speed-up with 10x less memory consumption using their proposal against dockerised containers. However, they also only used the one compiler for the generation of WebAssembly modules in their evaluations.

Jeong et al. (2019) proposed an edge computing framework for offloading JavaScript and WebAssembly state between mobile devices, edge computing devices and cloud services. This system saves the state of JavaScript workers, objects and WebAssembly functions, and enables the transfer and restoration of these between devices, resulting in an 8.4x speedup compared with only offloading pure JavaScript code. All the WebAssembly files used in their experiments were generated from C++ code using the Emscripten compiler. Once again, this highlights the gap of knowing how WebAssembly from other compilers would have performed in these experiments.

WebAssembly was proposed in (Long et al., 2021) as a lightweight alternative for serverless Function as a Service environments, that could enable fine grained resource consumption at runtime, as WebAssembly functions can be started and stopped on demand. They concluded that WebAssembly VMs outperform Docker containers and have a smaller footprint. However, only the Emscripten WebAssembly compiler was used to generate the WebAssembly bytecode for their evaluations.

Hall and Ramachandran (2019) also proposed WebAssembly as an alternative to traditional Dockerised containers for serverless functions in edge computing. They found that WebAssembly can provide many of the same isolation requirements, such as memory safety and security, while eliminating the cold start time penalty that

---

[4]https://llvm.org/

traditional serverless containers incur. Once again, the WebAssembly modules used in their experiments were all generated from C++ code compiled into WebAssembly using the Emscripten compiler.

Koren (2021) presented a proof-of-concept for a standalone WebAssembly development environment, capable of running on the edge and in the cloud. It features a built-in web-server and a browser-based Integrated Development Environment (IDE) as an alternative to running containerized microservices on low powered IoT devices with limited capabilities. Their evaluations were also done using WebAssembly modules generated from a single compiler, this time, the AssemblyScript compiler.

Tiwary et al. (2020) proposed an alternative to container-based serverless computing, addressing their cold-start problems and complicated architectures, while providing stateful memory and multi-tenancy isolation. Their proposal suggests WebAssembly to be executed directly on ring 0 with the serverless functions placed as close to the data as possible, providing CPU memory and filesystem isolation that is required for multi-tenancy in cloud computing. Their proposal was developed in Rust and compiled into WebAssembly using the Rust WebAssembly compiler, therefore once again, performance of WebAssembly is based on the code generated from a single WebAssembly compiler.

(Zheng et al., 2020) evaluate the performance of Ethereum blockchain smart contracts using WebAssembly. They implement 12 benchmarks in Rust and compile to WebAssembly. They also use the emerging compiler SOLL[5] that generates WebAssembly bytecode from Solidity[6]. Their experiments compare the performance of the WebAssembly driven blockchain against the EVM (Ethereum Virtual Machine) implementation, finding that WebAssembly driven blockchain does not perform as well as the EVM. However, they did not use the full range of available WebAssembly compilers by implementing their benchmarks in other programming languages, which perhaps might have produced different results for the evaluation of WebAssembly.

FAUST is a domain specific functional programming language used in audio pro-

---

[5]https://github.com/second-state/SOLL

[6]https://soliditylang.org/

cessing and sound synthesis for applications such as synthesizers, musical instruments and sound effects used in concerts, artistic productions, education and research. Letz et al. (2018) propose an alternative by compiling FAUST into WebAssembly and performing benchmarks to compare the performance of WebAssembly with native versions. They found that WebAssembly ran up to 66% slower than native C++ versions. Interestingly, they question the code quality that will eventually be generated by multiple WebAssembly compilers. However, their experiments only involved the WebAssembly modules generated by the Emscripten compiler.

Taheri (2018) discuss how computer vision on the web suffers from performance issues due to limitations in JavaScript. They developed a WebAssembly implementation of the OpenCV computer-vision library, which was initially implemented in C++. The library was compiled into WebAssembly using Emscripten and the performance was evaluated against the original implementation. Their findings were that WebAssembly did indeed run at close to native speeds, however only one WebAssembly compiler was used to generate the WASM bytecode.

This pattern of only using one WebAssembly compiler to benchmark the performance of WebAssembly code repeats itself through all the literature we reviewed. Murphy et al. (2020) only used the Clang compiler to evaluate WebAssembly for serverless computing. Cabrera Arteaga et al. (2020) proposed a toolchain for the superoptimization of WebAssembly binary code to improve execution time and overall file size. However, their experiments are limited to just C code using the Emscripten and LLVM compilers.

In this section we discussed some of the available literature that evaluated the performance of WebAssembly generated from programming languages such as C, C++, Rust and AssemblyScript using their respective compilers. However, none of these evaluations compared the code generated from these compilers against each other. As we highlight in this paper, each WebAssembly compiler generates different code even for simple programs, which produce significant differences in performance times.

Chapter 3

Experiment design and methodology

The entire source code and results of our experiments are publicly available on a GitHub Repository[1].

After installing the WebAssembly compilers for each of the selected languages and their required dependencies, this experiment can be fully reproduced by cloning the Git repository on a local VM and following the steps provided on the Github repository and in this paper.

3.0.1 Hypotheses

As an example, if we write a simple Fibonacci algorithm in C/C++, AssemblyScript and Rust, maintaining identical syntax between the languages, will they all have the same performance when compiled into WebAssembly and run on the same environment? In order to answer our research question using the scientific methods mentioned in chapter 1, we propose the following hypotheses:

---

[1]https://github.com/rayphelan/MastersProject2021

H0 - Null Hpyothesis:

> *There is no significant difference in the performance of WebAssembly modules, in terms of execution times, for the same program written in different programming languages and compiled into WebAssembly using their respective compilers.*

H1 - Alternate Hpyothesis:

> *There is a significant difference in the performance of WebAssembly modules, in terms of execution times, for the same program written in different programming languages and compiled into WebAssembly using their respective compilers.*

An important assumption of our experiments is that each of the algorithms across the selected programming languages is provided with the same input parameters and that the output of these algorithms is also the same for each language. This ensures that the results of the benchmarks are comparable, as highlighted in (Khan et al., 2015). For this reason, we have hardcoded the input parameters into each of the algorithms to ensure that each version of the algorithm in each language will have the same workload to operate on.

Furthermore, the objective of our experiments is not to benchmark the printing capabilities of each compiled WebAssembly algorithm, therefore only one single output is printed at the end of each algorithm execution for the purpose of visually confirming that the algorithm executed correctly.

3.1   Experiment Set up

In the following sections will provide details on the VM set up, Software Installation, Programming Languages and Algorithms used to perform our experiments. We also provide details of how the benchmarks were executed and how the results were gathered.

3.1.1 VM and Host Environment

The VM was installed on top of a Windows desktop computer. The details of both are listed below.

Table 3.1 show the details of the VM used in the experiments.

Operating System	Ubuntu 20.10 (64-bit)
RAM	12075 MB

Table 3.1: VM

Table 3.2 show the details of the environment used for hosting the VM.

Operating System	Windows 10 Enterprise (64-bit)
RAM	32 GB
Processor	Intel(R) Core(TM) i7-10610U CPU @ 1.80GHz 2.30 GHz

Table 3.2: Host Environment

3.1.2 Software and Compilers

The following section provides details on the software that was required. This includes details on the programming languages and the WebAssembly compilers used in the experiments.

Emscripten

Emscripten[2] is the compiler used to generate the WebAssembly files from C/C++ code. Installing Emscripten has a number of prerequisites, such as GIT[3], NodeJS (Version 15.14.0)[4] and Python3[5] and the GCC Compiler[6] for C and C++.

---

[2]https://emscripten.org/
[3]https://git-scm.com/
[4]https://nodejs.org/en/
[5]https://www.python.org/downloads/
[6]https://gcc.gnu.org/

Rust

Rust was downloaded and installed from here[7] and the documentation for setting up the Rust to WebAssemby compiler, wasm-pack[8] was found at here[9].

AssemblyScript

Documentation for setting up the AssemblyScript WebAssembly compiler can be found here[10].

RStudio

RStudio[11] was used to perform the statistical analysis of our results.

3.2   Algorithms

The algorithms used in our experiments were first sourced on the internet as C/C++ programs. Then the equivalent of these algorithms were written in AssemblyScript and Rust programming languages. We created a directory for each of the selected programming languages. Inside each of these directories, we created a directory for each of the selected algorithms. The algorithms were coded in their respective folder for their respective programming language. The algorithms were then compiled to WebAssembly using their respective compilers. The selection of algorithms used includes searching, sorting and numerical computing algorithms.

In our experiments, some of the algorithms were given an array of 100,000 elements as their input. A JavaScript program was created which generated 3 versions of an array with 100,000 elements. Each element is a unique floating point number with 7 decimal places. The first version is a sorted array with 100,000 elements. The second version is the same array in reserve order. The third version is the array with the

---

[7]https://www.rust-lang.org/
[8]https://rustwasm.github.io/wasm-pack/installer/
[9]https://rustwasm.github.io/docs/book/
[10]https://www.assemblyscript.org/
[11]https://rstudio.com/

elements in random positions. The respective array for each experiment was copied directly into the algorithms as their input parameters and the same array was provided to all algorithms which required an array as the input parameter.

The selection of algorithms used in this research were influenced by other research papers that previously used these algorithms. The NQueens algorithm has been used in papers such as (Herrera et al., 2018) and (Khan et al., 2015). Recursive functions, such as the Fibonacci algorithm, are used in the following benchmarking paper (van Eekelen et al., 2019). Reiser and Bläser (2017) used a number of algorithms for their benchmarks, including the Prime Numbers, Merge Sort and Fibonacci algorithms. The Bubble Sort, Merge Sort, Fibonacci, NQueens, Prime Number were also used in previous versions of the Ostrich Benchmarking Suite[12].

3.2.1   Numerical Computing Algorithms

The numerical computing algorithms used in our experiments are NQueens, Primary Number and Fibonacci Number algorithms.

NQueens

The 8 Queens Puzzle[13] was originally published in 1848. The objective is to place 8 Chess Queens on an 8x8 chess board in such a manor that no Queens are in direct threat from each other. Given the complexity of this puzzle when used with larger board sizes, it has been used many times as a software algorithm for benchmarking, known as the NQueens back-tracking algorithm.

The NQueens algorithm is part the Ostrich Benchmark Suite[14] developed at the Sable Lab at McGill University. It has been a popular research topic in computer science research for many years (Alhassan, 2019; Ayub et al., 2018; Guldal et al., 2016).

In our experiments, we implemented two versions of the NQueens algorithm. The

---

[12]https://github.com/Sable/Ostrich2

[13]https://en.wikipedia.org/w/index.php?title=Eight$_queens_puzzleoldid$ = 1022941393

[14]http://www.sable.mcgill.ca/mclab/projects/ostrich/

first using a board of 24x24, the second using a board of 27x27. The reason we chose the 24x24 size was that it provided enough complexity between all the selected programming languages in order to return a benchmarked time other than microseconds. The reason we did not go beyond the 27x27 size is because AssemblyScript was already taking considerably longer than the other programming languages the larger the board size became. The algorithm was implemented in all of the selected programming languages keeping the same code structure as much as possible. The output of the algorithms was the printed solution of the board.

The C/C++ version of this algorithm was sourced online[15]. The NQueens algorithm has a complexity of $O(n^n)$.

Prime Number

An important topic in computer science and cryptography is Number Theory, where the Prime Number calculation plays an important role (Elhakeem Abd Elnaby & El-Baz, 2021). A Prime Number is a number that has only two factors, 1 and the number itself. In this experiment, the Prime Numbers algorithm will calculate every prime number from 2 until 100,000. The output will print the 100,000th prime number. The C/C++ version of this algorithm was sourced online[16] and the code for the other selected languages was written from scratch. The Prime Numbers algorithm has a complexity of $O(\sqrt{}(n))$.

Fibonacci

The Fibonacci sequence is another well known computer algorithm used for benchmarking and was recently used in (Mendki, 2020) for benchmarking WebAssembly performance. It takes a number as the input parameter, and adds the previous two numbers together to produce the next Fibonacci number sequence of the given input. For our experiment, we used the recursive method to calculate the Fibonacci sequence. The algorithm will return the 45th Fibonacci sequence number. The reason

---

[15]https://www.geeksforgeeks.org/c-program-for-n-queen-problem-backtracking-3/

[16]https://www.csinfo360.com/2020/01/write-program-to-find-nth-prime-number.html

that we limited our experiment to calculating only the 45th number and not higher, is that it already provided enough complexity in the selected programming languages to perform our evaluations, with the execution times benchmarking between 5 and 10 seconds between the languages. The Fibonacci algorithm has a complexity of $O(\log n)$.

3.2.2 Sorting Algorithms

Sorting plays a crucial role in many algorithms (Abhay et al., 2019). The sorting algorithms used in this experiment are Bubble Sort, Heap Sort, Merge Sort, Selection Sort and Shell Sort. Each algorithm is given the same array of 100,000 elements as the input. Each element consists of a unique floating point number between 0 and 99,999 with 7 decimal places. The input array is sorted in reverse order, providing the algorithm with the worst case scenario, and the output of the algorithms will produce the same array in sorted order.

Bubble Sort

The Bubble Sort algorithm compares two adjacent elements and swaps them around if they are not in the required order. The C/C++ version of this algorithm was sourced online[17] and the algorithms were written in the other languages from scratch. The Bubble Sort algorithm has a complexity of $O(n2)$.

Heap Sort

The Heap Sort algorithm is generally slower than other sorting algorithms and therefore not commonly used, however we wanted to include this inefficient sorting algorithm in our experiments as we believe it still provides interesting insights. The C/C++ version of this algorithm was sourced online[18] and the algorithms were written in the other languages from scratch. The Heap Sort algorithm has a complexity of $O(n \log n)$.

---

[17]https://https://www.programiz.com/dsa/bubble-sort
[18]https://www.programiz.com/dsa/heap-sort

Merge Sort

The Merge Sort algorithm follows the principle of Divide and Conquer. It works by dividing the problem into multiple sub-problems, solving each of the smaller sub-problems, and merging the results back together to generate the output. The C/C++ versions of this algorithm was sourced online[19] and the algorithms were written in the other languages from scratch. The Merge Sort algorithm has a complexity of $O(n * \log n)$.

Selection Sort

The Selection Sort algorithm selects the smallest element from an array and places it at the beginning. It repeats this process until the results represent a sorted array. The C/C++ version of this algorithm was sourced online[20] and the algorithms were written in the other languages from scratch. The Selection Sort algorithm has a complexity of $O(n2)$.

Shell Sort

The Shell Sort algorithm sorts elements at different intervals, starting from the elements that are furthest apart from each other and repeating the process until all the elements are in a sorted order. The C/C++ version of this algorithm was sourced online[21] and the algorithms were written in the other languages from scratch. The Shell Sort algorithm has a complexity of $O(n2)$.

3.2.3 Searching Algorithms

The most commonly known searching algorithms are the Binary Search and Linear Search algorithms (Jacob et al., 2018), therefore we have implemented these algorithms in our experiments. The input for these algorithms consist of an array of 100,000 elements of unique floating point numbers between 0 and 99,999 with 7 decimal places.

---

[19]https://www.programiz.com/dsa/merge-sort

[20]https://www.programiz.com/dsa/selection-sort

[21]https://www.programiz.com/dsa/shell-sort

The algorithm will run a loop from 0 to 99,999 searching for the number of each iteration within the input array. The output will be the index of the array where the searched element was found on the last iteration.

Binary Search

The Binary Search algorithm assumes that the given array will in a sorted order. Given a sorted array, the algorithm divides the array in half by selecting the middle element of the array. It determines if the number being searched for is greater or less than the value of the middle element. This means that one half of the array can be ignored as the number being searched for must be in the other half. This process is repeated recursively until the number being searched for is found. The C/C++ code for this algorithm was sourced online[22] and the algorithms were written in the other languages from scratch. The Binary Search has a complexity of $O(\log n)$.

Linear Search

The Linear Search algorithm does not require a sorted array. In this experiment, the Linear Search algorithm was given an array of 100,000 elements of unique floating point numbers between 0 and 99,999 with 7 decimal places and the array has been given a random ordering. The algorithm performs a simple search on every element in the array, starting from the first element and finishing when the number being searched for is found, which could go all the way to the last element. The C/C++ code for this algorithm was sourced online[23] and the algorithms were written in the other languages from scratch. The Linear Search has a complexity of $O(n)$.

3.3 Compiling to WebAssembly

This section will provide details on how each WebAssembly compiler was used. Each of the compilers come with optimization level options. Some optimizations can be

---

[22]https://www.programiz.com/dsa/binary-search
[23]https://www.programiz.com/dsa/linear-search

made which will reduce the overall code size. Other optimizations can be made which increase performance in execution time. In our experiments, we chose to optimize for execution time speed rather than code size.

3.3.1   C/C++ to WebAssembly

Emscripten is used to compile C and C++ into WebAssembly. It takes in a C program as an argument and generates a WebAsssembly file and a JavaScript glue file. Emscripten allows for different levels of optimization, ranging from no optimization to highly optimized. For our experiments, we chose the optimization level -O2 as this provided a combination of well optimized code with fast compile times.

The optimization levels that Emscripten provides are listed in table 3.3.

Optimization Level	Description
-O0	No optimization
-O1	Low optimization, shorter compile time
-O2	Well optimized build
-O3	Highly optimized, longer compile time
-Os	Highly optimized, reduced code size, longer compile time

Table 3.3: Emscripten Optimization Levels

Emscripten also takes an arguement to modularize the code. This means that the compiler will generate the JavaScript glue file in such a way that allows it to be imported as a JavaScript module. The compiler puts all of the required JavaScript into a factory function which can be called on to create an instance of the WebAssembly module. The following code in listing 3.1 shows how the WebAssembly file and the JavaScript glue file were generated using Emscripten for the BinarySearch algorithm.

Listing 3.1: Compiling BinarySearch from C to WebAssembly

```
1  emcc binarySearch.c -o binarySearch.js -O2
2      -s MODULARIZE -s "EXPORTED_FUNCTIONS=['_binarySearch']"
```

3.3.2   AssemblyScript to WebAssembly

AssemblyScript was designed specifically to target WebAssembly while offering a familiar syntax for TypeScript and JavaScript developers. It uses Node Package Manager[24] (NPM) as the installer. We created a new NPM application for each algorithm. Instructions for setting up new NPM applications are detailed on our Github Repository. The AssemblyScript compiler also comes with options for optimization. For our experiments we used the optimization level 3, for speed rather than code size, like we did with the Emscripten optimization levels. However, these optimization levels are actually applied by default, but can be altered in the configuration files of the compiler. The following code in listing 3.2 show how the WebAssembly and JavaScript files were generated using AssemblyScript.

Listing 3.2: AssemblyScript Optimization

```
1   npm run asbuild
```

3.3.3   Rust to WebAssembly

To compile Rust to WebAssembly, the wasm-pack compiler was used. Instructions for creating new wasm-pack applications can be found on our Github Repository. We created a new wasm-pack application for each of our selected algorithms. Once inside the directory of the wasm-pack application created for each algorithm, the WebAssembly modules can be compiled.

We passed an option to the compiler to note that we were targetting the NodeJS environment. This customizes the JavaScript glue file so it can be easily run in NodeJS. Wasm-pack allows us to specify whether or not the build should be optimized for production, using the –release keyword. However if no keyword is provided, the –release profile is automatically used. This was the optimization levels used in our experiments. The following code in listing 3.3 shows how the WebAssembly and JavaScript files were generated for Rust.

---

[24]https://www.npmjs.com/

Listing 3.3: Rust to WebAssembly

```
1  wasm—pack build ——target nodejs
```

3.4 Benchmarking and Gathering Results

WebAssembly modules are most commonly loaded via its JavaScript API (Hall & Ramachandran, 2019). The WebAssembly compilers used in our experiments generate the boilerplate JavaScript glue which can then be be converted into a JavaScript module, allowing it to be conveniently imported into other JavaScript functions. Inspired by the Ostrich Benchmark Suite which created a python runner script for executing the benchmarks and recording the results, we created our own JavaScript runner which runs our experiments in NodeJS with the desired number of iterations and records the final results onto a CSV (Comma Separated Values) file for post processing.

Each compiler generates the WASM file and the JavaScript glue for loading and interacting with the WASM file. For each algorithm in each language, we created a new JavaScript runner which loads the JavaScript Glue file. The JavaScript runner was manually invoked using NodeJS command line for each algorithm in each programming language. Once the WebAssembly module has been loaded, the JavaScript runner executes a loop which runs the main function in the corresponding WebAssembly module. The number of iterations in the loop depends on how many times we want to perform the experiment. After each iteration, the execution time in nanoseconds is measured and converted into seconds and milliseconds before being saved to an array of results. After the final iteration, the array is saved into a CSV file. Note that the time for loading the WebAssembly module is not being measured, since we are only interested in the execution time.

Herrera et al. (2018) suggested running each experiment 30 times, (Mendki, 2020) ran their experiments 100 times, while the JetStream2 Benchmark[25] suggested running the experiments 120 times. Initially we only ran our experiments 30 times each. However, upon visual inspection of the distribution of the results using histograms, we

---

[25]https://https://browserbench.org/JetStream/in-depth.html

found that our data was not normalised. We then proceeded to run the experiments 120 times each. However, even after running the experiments 120 times each, the results mainly did not have a normal distribution. This would influence our choice in statistical tools during our evaluation of the results.

The code in listing 3.4 represents the JavaScript runner for the Fibonacci algorithm in C. Figure 3.1 shows a sequence diagram of the JavaScript Runner file we created.

Listing 3.4: JavaScript runner for Fibonacci in C WebAssembly

```
1   const wasmModule = require('./fibonacci.js');
2   const fs = require('fs');
3   const results = [];
4   const iterations = 120;
5   wasmModule().then((instance) => {
6       for (n = 1; n <= iterations; n++) {
7
8           // Start Timer
9           const start = process.hrtime();
10
11          // Run WASM
12          const wasm = instance._fibonacci();
13
14          // End Timer
15          const diff = process.hrtime(start);
16
17          const result = (diff[0] * 1e9 + diff[1])/1000000000;
18          results.push(result);
19          console.log(wasm);
20      }
21      const csv = results.join('\n');
22
23      // Write File
24      fs.writeFile('results.csv', csv, function (err) {
25          if (err) return console.log(err);
26          console.log('Filesaved');
27      });
28  });
```
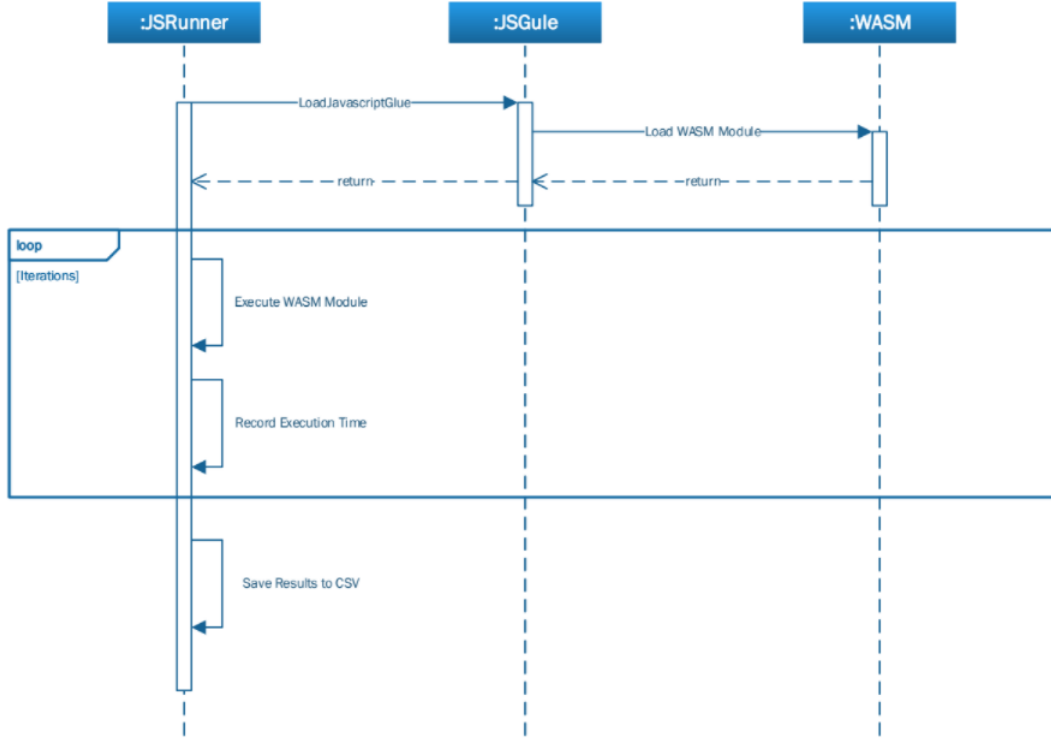
Figure 3.1: Experiment JavaScript Runner

When all of the experiments had been completed, we created Excel files for each language (C, C++, AssemblyScript, Rust). Each column in the spreadsheet represents the algorithm being used in the experiment. Each row represents the execution time of the algorithm for the number of iterations performed. The Excel files were then saved as CSV files for further processing in RStudio to perform the statistical analysis.

Figure 3.2 shows an example of how the results were stored in Excel once the experiments have been finished.

	A	B	C	D	E	
1	BinarySearch	BubbleSort	Fibonacci	HeapSort	LinearSearch	MergeS
2	0.024169105	7.33237846	6.096046528	0.022217607	4.882848929	
3	0.032136363	8.831246958	6.495604445	0.018852237	5.914289661	
4	0.035062501	8.317894254	7.213856816	0.015125697	5.839858018	
5	0.02848115	8.282457763	6.78204264	0.025129889	5.806848937	

Figure 3.2: Gathering Results

3.5   Analysis of Results

RStudio was used to perform the statistical analysis of the results. RStudio provides a convenient way of running statistical tests for normality and significance. It also provides functionality to plot graphs for visual interpretation of the results.

Using the previously generated CSV files with the execution times of the algorithms in each language, these were imported into RStudio. The required RStudio libraries were also included. Figure 3.3 shows how the libraries were included and how the CSV files were imported. The entire source code used in RStudio is included in the Appendix section of this paper and on our Github repository.

```
 1
 2  library("dplyr")
 3  library("ggpubr")
 4
 5
 6 ~ ############################################################################
 7  # Load Languages
 8  AssemblyScript <- read.csv('C:/TuDublin/Masters Project/Git Clone/MastersProject2021/FinalResultsPerLanguage/CSV/AssemblyScript.csv')
 9  C <- read.csv('C:/TuDublin/Masters Project/Git Clone/MastersProject2021/FinalResultsPerLanguage/CSV/C.csv')
10  CPP <- read.csv('C:/TuDublin/Masters Project/Git Clone/MastersProject2021/FinalResultsPerLanguage/CSV/CPP.csv')
11  Rust <- read.csv('C:/TuDublin/Masters Project/Git Clone/MastersProject2021/FinalResultsPerLanguage/CSV/Rust.csv')
12
```

Figure 3.3: Importing Data into RStudio

3.5.1   Testing for Normal Distribution

It is important that we understand the distribution of data in our results, as this will determine how we test for significant differences in the execution times, either with parametric tests for a normal distribution, or non-parametric tests for a non-normal distribution.

We ran the Shapiro Wilk[26] test on each programming language result file for each algorithm. If the results from the test return a value of 0.05 or more, then it means our data has a normal distribution, otherwise it has a non-normal distribution.

In RStudio, running this test is quite a simple process. The following listing 3.5 shows and example of the R commands used to perform the Shapiro Wilk test on each algorithm for the AssemblyScript results. This process was repeated for each of the other programming languages, C, C++ and Rust.

---

[26]https://statistics.laerd.com/spss-tutorials/testing-for-normality-using-spss-statistics.php

Listing 3.5: Shapiro Wilk Test in RStudio

```
1   shapiro.test(AssemblyScript$BinarySearch)
2   shapiro.test(AssemblyScript$BubbleSort)
3   shapiro.test(AssemblyScript$Fibonacci)
4   shapiro.test(AssemblyScript$HeapSort)
5   shapiro.test(AssemblyScript$LinearSearch)
6   shapiro.test(AssemblyScript$MergeSort)
7   shapiro.test(AssemblyScript$NQueen24)
8   shapiro.test(AssemblyScript$NQueen27)
9   shapiro.test(AssemblyScript$PrimeNumber)
10  shapiro.test(AssemblyScript$SelectionSort)
11  shapiro.test(AssemblyScript$ShellSort)
```

The results of the Shapiro Wilk tests for normality showing the p-value confirms that our results were not normally distributed. The results can be seen in table 3.4. For the purpose of facilitating the visualisation of the data, we will abbreviate AssemblyScript to AS.

	AS	C	C++	Rust
BinarySearch	9.081e-08	2.251e-09	4.471e-08	0.4335
BubbleSort	0.000131	2.95e-07	5.32e-09	8.002e-07
Fibonacci	1.921e-06	1.394e-08	7.82e-09	2.501e-07
HeapSort	2.93e-06	5.966e-06	4.656e-06	3.685e-07
LinearSearch	1.97e-09	9.529e-06	8.256e-05	< 2.2e-16
MergeSort	1.878e-06	7.426e-06	2.564e-09	< 2.2e-16
NQueen24	0.0005585	0.0002786	< 2.2e-16	0.01424
NQueen27	0.0009603	1.927e-13	2.8e-11	8.382e-12
PrimeNumber	1.986e-06	8.051e-09	1.183e-12	2.399e-08
SelectionSort	1.995e-13	2.365e-06	1.451e-06	4.591e-08
ShellSort	1.713e-07	9.023e-12	1.832e-06	0.04334

Table 3.4: Shapiro Wilk Test for Normality

The Shapiro Wilk test showed that our data was not normally distributed. The

Central Limit Theorem[27] states that if we have sufficiently large data samples, then the distribution will be approximately normal and parametric tests can be used. However, if the sample size is not sufficiently large and the distribution of data is not normal, then non-parametric tests can be used to determine the significant difference between the samples.

At this stage, we had to choose between running the experiments again a greater number of times until we achieved a normal distribution, so we could use parametric tests on our results, or keeping them as they are and using non-parametric tests on the results. Given the time constraints on completing this paper, we decided to use the results we had already obtained and apply non-parametric tests.

As a visual aid, we also generated some graphs to confirm the tests. We created visual graphs in both RStudio and Tableau Public[28]. The following figure 3.4 created in Tableau Public shows the distribution of the results of the Fibonacci algorithm in the selected languages. The entire collection of graphs can be found in the Appendix of this paper.
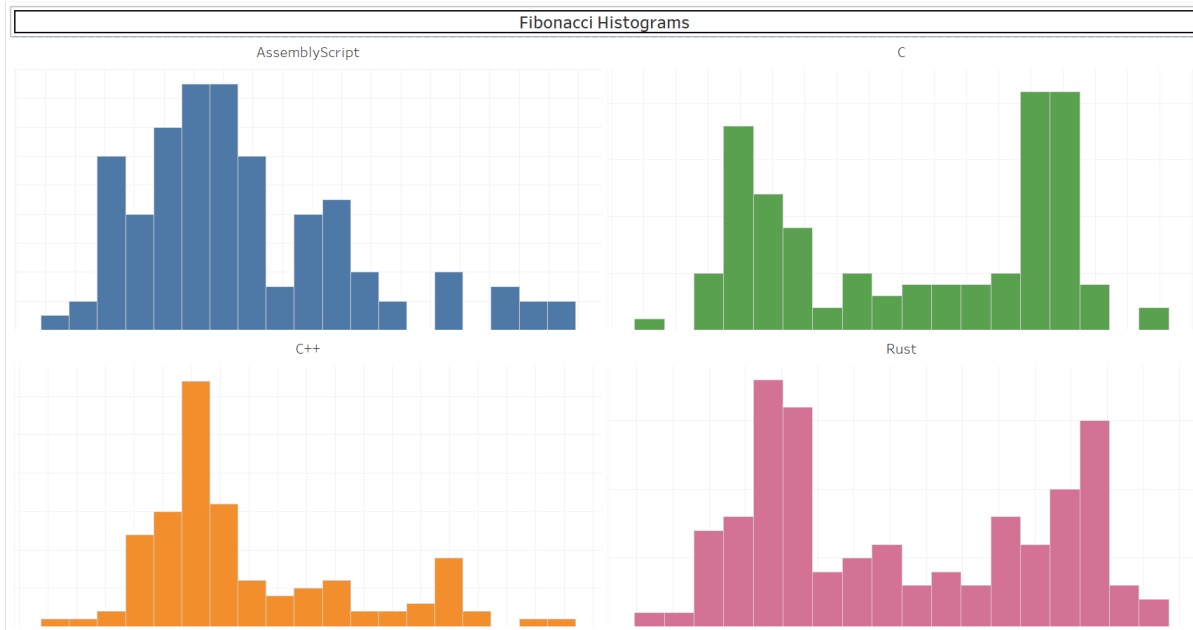
Figure 3.4: Fibonacci Histograms

The following graphs created in RStudio, shown in figure 3.5 represent the distribution of the results for each algorithm in AssemblyScript. The entire collection of these graphs can be found in the Appendix section of this paper.

Figure 3.5: AssemblyScript Distribtion

(a) BinarySearch

(b) BubbleSort

(c) Fibonacci

(d) HeapSort

(e) LinearSearch

(f) MergeSort

(g) NQueen24

(h) NQueen27

(i) PrimaryNumber

(j) SelectionSort

(k) ShellSort
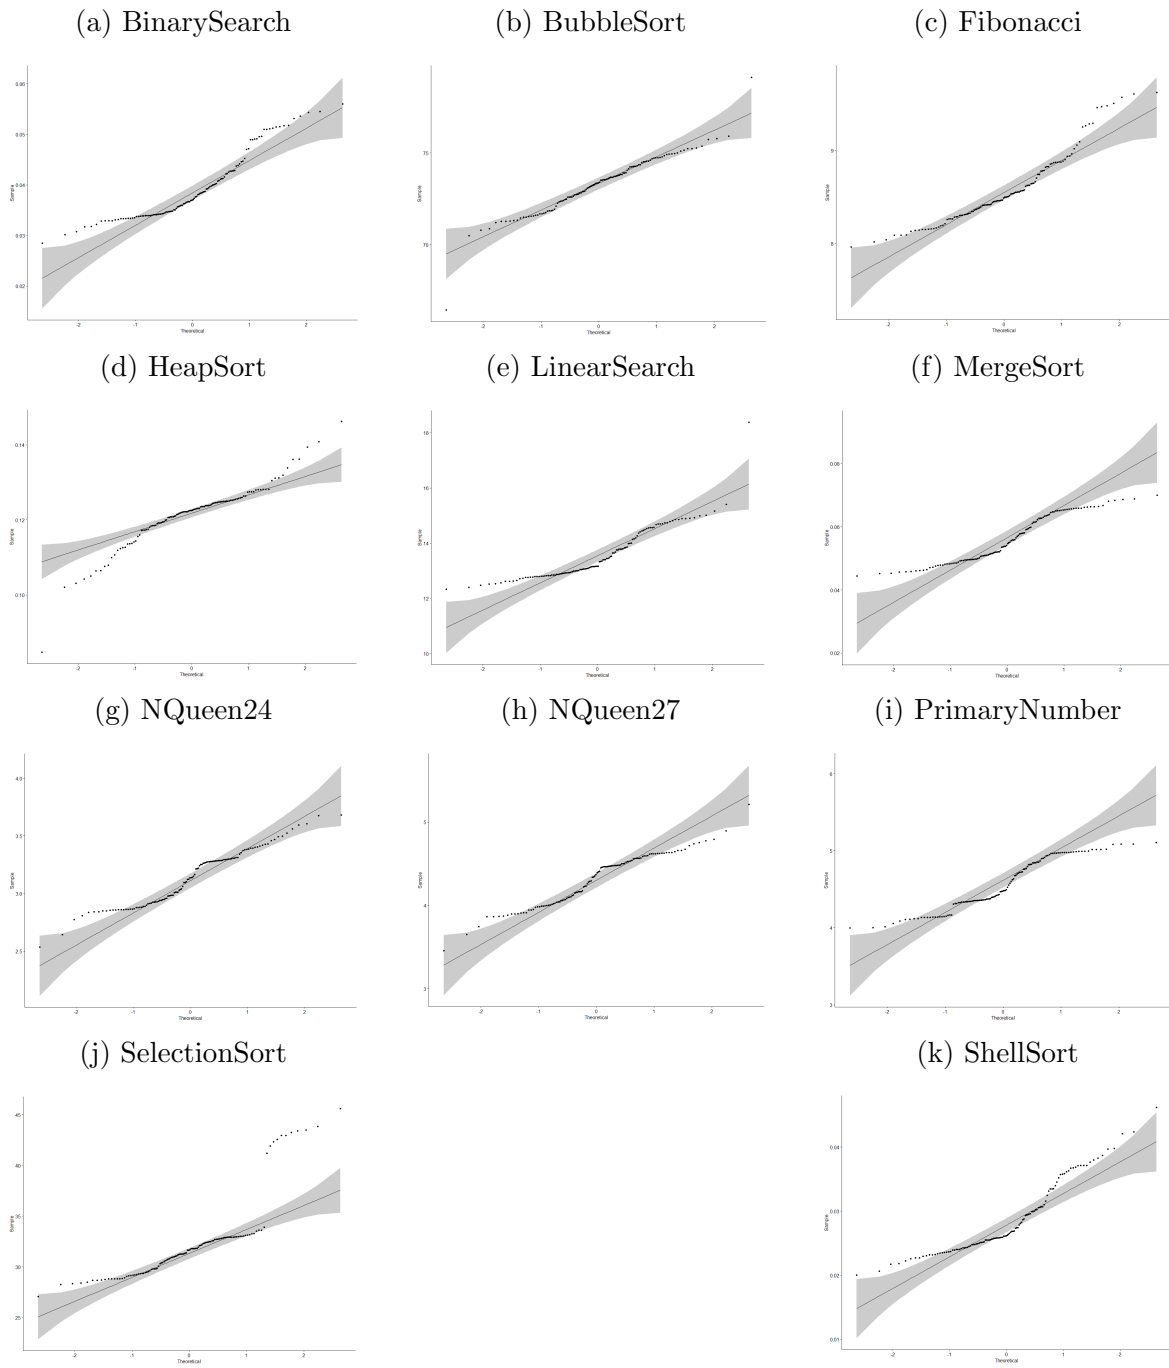
3.5.2   Testing for Significant Difference

In order to accept or reject the null hypothesis, we need to determine whether there is a significant difference in the execution times of the algorithms when compared to each other. We grouped the languages into pairs and ran a non-parametric test on them to obtain the p-value. Having a p-value of 0.05 or less means that we can reject the null hypothesis as there is significant difference between the two samples.

The non-parametric test we chose is the Mann-Whitney U Test, which is used to compare exactly two independent samples and the data is not normally distributed. We ran the test on each algorithm for AssemblyScript vs C, AssemblyScript vs C++, AssemblyScript vs Rust, C vs C++, C vs Rust and C++ vs Rust.

The following code in listing 3.6 shows how the Mann-Whitney U Test was run on the BinarySearch algorithm for AssembyScript vs C, AssemblyScript vs C++ and AssemblyScript vs Rust. The entire RStudio code can be found in the Appendix of this paper and on our Github repository.

Listing 3.6: Mann-Whitney U Test in RStudio

```
1   wilcox.test(AssemblyScript$BinarySearch,C$BinarySearch)
2   wilcox.test(AssemblyScript$BinarySearch,CPP$BinarySearch)
3   wilcox.test(AssemblyScript$BinarySearch,Rust$BinarySearch)
```

The values represented in following the results are the p-value. In order to facilitate the visualisataion of the results, we have divided the results into two tables. The first table 3.5 shows the results of AssemblyScript vs C, AssemblyScript vs C++ and AssemblyScript vs Rust. The second table 3.6 shows the results of C vs C++, C vs Rust and C++ vs Rust.

	AS vs C	AS vs C++	AS vs Rust
BinarySearch	< 2.2e-16	< 2.2e-16	< 2.2e-16
BubbleSort	< 2.2e-16	< 2.2e-16	< 2.2e-16
Fibonacci	< 2.2e-16	< 2.2e-16	< 2.2e-16
HeapSort	< 2.2e-16	< 2.2e-16	< 2.2e-16
LinearSearch	< 2.2e-16	< 2.2e-16	< 2.2e-16
MergeSort	< 2.2e-16	< 2.2e-16	< 2.2e-16
NQueen24	< 2.2e-16	< 2.2e-16	< 2.2e-16
NQueen27	< 2.2e-16	< 2.2e-16	< 2.2e-16
PrimeNumber	< 2.2e-16	< 2.2e-16	< 2.2e-16
SelectionSort	< 2.2e-16	< 2.2e-16	< 2.2e-16
ShellSort	< 2.2e-16	< 2.2e-16	8.265e-06

Table 3.5: Mann-Whitney U Test

	C vs C++	C vs Rust	C++ vs Rust
BinarySearch	0.1421	2.427e-09	7.219e-12
BubbleSort	3.225e-13	< 2.2e-16	< 2.2e-16
Fibonacci	0.07379	< 2.2e-16	4.606e-16
HeapSort	0.6791	1.119e-13	2.901e-14
LinearSearch	0.001119	< 2.2e-16	< 2.2e-16
MergeSort	0.6215	< 2.2e-16	< 2.2e-16
NQueen24	< 2.2e-16	0.01026	< 2.2e-16
NQueen27	< 2.2e-16	< 2.2e-16	< 2.2e-16
PrimeNumber	2.354e-06	< 2.2e-16	< 2.2e-16
SelectionSort	0.001779	< 2.2e-16	< 2.2e-16
ShellSort	0.4074	< 2.2e-16	< 2.2e-16

Table 3.6: Mann-Whitney U Test

Chapter 4

Results, evaluation and discussion

The main objective of our experiments was to answer the research question by either accepting or rejecting the null hypothesis.

Research Question:

> *Is there a significant difference in the performance of WebAssembly modules, in terms of execution times, for the same program written in different programming languages and compiled into WebAssembly using their respective compilers?*

H0 - Null Hpyothesis:

> *There is no significant difference in the performance of WebAssembly modules, in terms of execution times, for the same program written in different programming languages and compiled into WebAssembly using their respective compilers.*

H1 - Alternate Hpyothesis:

> *There is a significant difference in the performance of WebAssembly modules, in terms of execution times, for the same program written in different programming languages and compiled into WebAssembly using their respective compilers.*

4.1 Statistical Tests

We ran statistical tests on our paired data groups to obtain the p-values, which determines whether or not there is a significant difference between grouped pairs. A p-value of 0.05 or less indicates that there is indeed a significant difference, therefore we must reject the null hypothesis and accept the alternate hypothesis. This can be summarized by the following formula:

$$p \leq 0.05 = H1$$

On the other hand, having a p-value of higher than 0.05 indicates that there is no significant difference between the pair, and we must accept the null hypothesis and reject the alternate hypothesis. This can be summarized by the following formula:

$$p > 0.05 = H0$$

The statistical tests clearly showed significant differences in almost all of the experiments. One exception was while comparing C to C++, which on some experiments had a p-value greater than 0.05. However, this is not surprising, since both C and C++ are compiled to WebAssembly using the same compiler and the code is almost identical. Figure 4.1 highlights the results where the p-values were greater than 0.05, meaning there were no significant difference in those execution times. The results also show that for all other language pairs, there were significant differences in the execution times.

	AssemblyScript vs C	AssemblyScript vs CPP	AssemblyScript vs Rust	C vs CPP	C vs Rust	CPP vs Rust
BinarySearch	2.20E-16	2.20E-16	2.20E-16	0.1421	2.43E-09	7.22E-12
LinearSearch	2.20E-16	2.20E-16	2.20E-16	0.001119	2.20E-16	2.20E-16
NQueen24	2.20E-16	2.20E-16	2.20E-16	2.20E-16	0.01026	2.20E-16
NQueen27	2.20E-16	2.20E-16	2.20E-16	2.20E-16	2.20E-16	2.20E-16
Fibonacci	2.20E-16	2.20E-16	2.20E-16	0.07379	2.20E-16	4.61E-16
PrimeNumber	2.20E-16	2.20E-16	2.20E-16	2.35E-06	2.20E-16	2.20E-16
HeapSort	2.20E-16	2.20E-16	2.20E-16	0.6791	1.12E-13	2.90E-14
MergeSort	2.20E-16	2.20E-16	2.20E-16	0.6215	2.20E-16	2.20E-16
ShellSort	2.20E-16	2.20E-16	8.27E-06	0.4074	2.20E-16	2.20E-16
SelectionSort	2.20E-16	2.20E-16	2.20E-16	0.001779	2.20E-16	2.20E-16
BubbleSort	2.20E-16	2.20E-16	2.20E-16	3.23E-13	2.20E-16	2.20E-16

Figure 4.1: P-Values

4.2 Visual Representations

In this section, we will provide a detailed evaluation of the results for each algorithm. In order to provide a visual representation of the execution times, Tableau Public was used to produce line-charts of the execution time of each language compiled into WebAssembly for each algorithm. The visualizations clearly show that there is a significant difference in execution time between the algorithms. For visual inspection of the WebAssembly code, we used the Wabt toolkit[1] to transform the WASM files into the human readable WAT format.

4.2.1 Numerical Computing Algorithms

The Numerical Computing algorithms used in our experiments are pure mathematical functions and did not require the input array of 100,000 elements.

Fibonacci

The figure 4.2 shows the execution times for the Fibonacci algorithm. C and C++ are consistently running with similar execution times. If fact the p-values for the Fibonacci algorithm for C and C++ showed that there was no significant difference in their execution times. However, Rust appears to run faster than C and C++, while it can be visually observed that AssemblyScript performed the slowest of all the other languages.

[1]https://github.com/WebAssembly/wabt

Figure 4.2: Fibonacci

In all of our selected programming languages, the code structure for the Fibonacci sequence have the same syntax. This can be confirmed in the following listings. As mentioned previously, we implemented the Fibonacci algorithm using the recursive method. The algorithms will produce an output with the 45th Fibonacci sequence number.

The following listing 4.1 shows the Fibonacci algorithm in AssemblyScript.

Listing 4.1: Fibonacci in AssemblyScript

```
1  function fib(n: i32): i32 {
2    return n <= 1 ? 1 : fib(n−1) + fib(n−2);
3  }
4  export function fibonacci(): i32 {
5    const n: i32 = 45;
6    const result: i32 = fib(n);
7    return result;
8  }
```

The following listing 4.2 shows the Fibonacci algorithm in C.

Listing 4.2: Fibonacci in C

```
1  int fib(n) {
```

```
2    if (n <= 1) return 1;
3    return fib(n − 1) + fib(n − 2);
4  }
5  int fibonacci() {
6    int n = 45;
7    int result = fib(n);
8    return result;
9  }
```

The following listing 4.3 shows the Fibonacci algorithm in C++.

Listing 4.3: Fibonacci in C++

```
1  extern "C" {
2    int fib(int n) {
3      if (n <= 1) return 1;
4      return fib(n − 1) + fib(n − 2);
5    }
6    int fibonacci() {
7      int n = 45;
8      int result = fib(n);
9      return result;
10   }
11 }
```

The following listing 4.4 shows the Fibonacci algorithm in Rust. One important difference with Rust is that we need to include the a library so it can interact with JavaScript.

Listing 4.4: Fibonacci in Rust

```
1  use wasm_bindgen::prelude::*;
2  #[wasm_bindgen]
3  extern "C" {
4      #[wasm_bindgen(js_namespace = console)]
5      fn log(s: u64);
6  }
7  fn fibonacci(n: i32) −> u64 {
8      if n <= 1 { return 1 }
9      return fibonacci(n−1) + fibonacci(n−2);
10 }
11 #[wasm_bindgen(start)]
12 pub fn main_js() {
13     log(fibonacci(45));
14 }
```

Given that the code syntax for the Fibonacci algorithm in each of the selected
languages is as consistent as possible by implementing the algorithm using the recur-
sive method, an interesting observation is that they produced considerably different
WebAssembly binaries when we performed the visual inspection of the WAT files.

The following listing 4.5 is the WAT file representation of the compiled Fibonacci
algorithm in AssemblyScript.

Listing 4.5: Fibonacci WAT format for AssemblyScript

```
1   ( module
2     ( type $none_=>_i32 ( func ( result i32 )))
3     ( type $i32_=>_i32 ( func ( param i32 ) ( result i32 )))
4     ( memory $0  0)
5     ( export "fibonacci" ( func $assembly/index/fibonacci ))
6     ( export "memory" ( memory $0 ))
7     ( func $assembly/index/fib ( param $0 i32 ) ( result i32 )
8       local . get $0
9       i32 . const 1
10      i32 . le_s
11      if ( result i32 )
12       i32 . const 1
13      else
14       local . get $0
15       i32 . const 1
16       i32 . sub
17       call $assembly/index/fib
18       local . get $0
19       i32 . const 2
20       i32 . sub
21       call $assembly/index/fib
22       i32 . add
23      end
24     )
25     ( func $assembly/index/fibonacci ( result i32 )
26      i32 . const 45
27      call $assembly/index/fib
28     )
29   )
```

The following code in listing 4.6 is the WAT file representation of the compiled
Fibonacci algorithm in C and C++. The WebAssembly file that was generated for
C++ was identical to the one generated for C, again this was to be expected as they

are both compiled using the same compiler. However, an interesting observation here is that there is considerably more code in the C and C++ version of the compiled WebAssembly when compared to AssemblyScript for the exact same algorithm.

Listing 4.6: Fibonacci WAT format for C and C++

```
1   (module
2     (type (;0;) (func (result i32)))
3     (type (;1;) (func (param i32) (result i32)))
4     (type (;2;) (func))
5     (type (;3;) (func (param i32)))
6     (func (;0;) (type 2)
7       nop)
8     (func (;1;) (type 1) (param i32) (result i32)
9       (local i32 i32)
10      i32.const 1
11      local.set 1
12      local.get 0
13      i32.const 2
14      i32.ge_s
15      if (result i32)   ;; label = @1
16        i32.const 0
17        local.set 1
18        loop   ;; label = @2
19          local.get 0
20          i32.const 1
21          i32.sub
22          call 1
23          local.get 1
24          i32.add
25          local.set 1
26          local.get 0
27          i32.const 3
28          i32.gt_s
29          local.set 2
30          local.get 0
31          i32.const 2
32          i32.sub
33          local.set 0
34          local.get 2
35          br_if 0 (;@2;)
36        end
37        local.get 1
38        i32.const 1
39        i32.add
```

```
40      else
41        local.get 1
42      end)
43    (func (;2;) (type 0) (result i32)
44      i32.const 45
45      call 1)
46    (func (;3;) (type 0) (result i32)
47      global.get 0)
48    (func (;4;) (type 3) (param i32)
49      local.get 0
50      global.set 0)
51    (func (;5;) (type 1) (param i32) (result i32)
52      global.get 0
53      local.get 0
54      i32.sub
55      i32.const −16
56      i32.and
57      local.tee 0
58      global.set 0
59      local.get 0)
60    (func (;6;) (type 0) (result i32)
61      i32.const 1024)
62    (table (;0;) 1 1 funcref)
63    (memory (;0;) 256 256)
64    (global (;0;) (mut i32) (i32.const 5243920))
65    (export "memory" (memory 0))
66    (export "__wasm_call_ctors" (func 0))
67    (export "fibonacci" (func 2))
68    (export "__errno_location" (func 6))
69    (export "stackSave" (func 3))
70    (export "stackRestore" (func 4))
71    (export "stackAlloc" (func 5))
72    (export "__indirect_function_table" (table 0)))
```

The following listing 4.7 is the WAT file representation of the compiled Fibonacci algorithm in Rust.

Listing 4.7: Fibonacci WAT format for Rust

```
1  (module
2    (type (;0;) (func))
3    (type (;1;) (func (param i32 i32)))
4    (type (;2;) (func (param i32 i32 i32)))
5    (type (;3;) (func (param i32) (result i32)))
6    (type (;4;) (func (param i64) (result i64)))
7    (import "__wbindgen_placeholder__" "__wbg_log_e2b7116aabd69db1" (func (;0;) (type 1)))
```

```
8    (func (;1;) (type 4) (param i64) (result i64)
9      (local i64)
10     i64.const 1
11     local.set 1
12     local.get 0
13     i64.const 2
14     i64.ge_u
15     if (result i64)   ;; label = @1
16       i64.const 0
17       local.set 1
18       loop   ;; label = @2
19         local.get 0
20         i64.const −1
21         i64.add
22         call 1
23         local.get 1
24         i64.add
25         local.set 1
26         local.get 0
27         i64.const −2
28         i64.add
29         local.tee 0
30         i64.const 1
31         i64.gt_u
32         br_if 0 (;@2;)
33       end
34       local.get 1
35       i64.const 1
36       i64.add
37     else
38       local.get 1
39     end)
40   (func (;2;) (type 2) (param i32 i32 i32)
41     (local i64)
42     local.get 0
43     local.get 1
44     i64.extend_i32_u
45     local.get 2
46     i64.extend_i32_u
47     i64.const 32
48     i64.shl
49     i64.or
50     call 1
51     local.tee 3
52     i64.store32
```

```
53        local.get 0
54        local.get 3
55        i64.const 32
56        i64.shr_u
57        i64.store32 offset=4)
58      (func (;3;) (type 0)
59        (local i64)
60        i64.const 40
61        call 1
62        local.tee 0
63        i32.wrap_i64
64        local.get 0
65        i64.const 32
66        i64.shr_u
67        i32.wrap_i64
68        call 0)
69      (func (;4;) (type 3) (param i32) (result i32)
70        local.get 0
71        global.get 0
72        i32.add
73        global.set 0
74        global.get 0)
75      (memory (;0;) 17)
76      (global (;0;) (mut i32) (i32.const 1048576))
77      (export "memory" (memory 0))
78      (export "fibonacci" (func 2))
79      (export "main_js" (func 3))
80      (export "__wbindgen_add_to_stack_pointer" (func 4))
81      (export "__wbindgen_start" (func 3)))
```

Based on our experiment and analysis of the Fibonacci algorithm, we can clearly see a significant difference in the performance of the WebAssembly modules from different programming languages, with the exception of C and C++ that produced identical WebAssembly files and had a p-value greater than 0.05. We can also see a clear difference of the WebAssembly files when converted into WAT format.

In order to not unnecessarily bloat the contents of this paper, the source code and WAT files for the remaining algorithms will not be included directly in the paper. However, they may be viewed on our Github repository.

PrimeNumber

The visualization in figure 4.3 shows that C and C++ perform with similar execution times, although their p-values were less than 0.05. We can also see that Rust performs slightly slower than C and C++. Both however generally performed the PrimeNumber algorithm in less than 1 second every time. On the other hand, we can clearly see that AssemblyScript took longer than all the other languages, taking between 4 to 5 seconds to complete.

Figure 4.3: PrimeNumber

We confirmed that the syntax for the PrimaryNumber algorithm is the same for each of the selected languages. However, the WebAssembly files generated were considerably different. We observed that the AssemblyScript produced 279 lines of code when the WAT extension file was inspected, the C and C++ versions produced only 36 lines of code, while the Rust version produced 7005 lines of code.

NQueen24 and NQueen27

The source code for the NQueens algorithms for each of the selected languages were kept syntactically the same as much as possible in our experiments. Upon inspection of the WAT files for each of the selected languages, we observed that AssemblyScript generated 3090 lines, C and C++ generated 3218 lines, while Rust generated 7622 lines.

The execution times of the NQueen24 and NQueen27 algorithms can be viewed in figure 4.4 and figure 4.5.

Figure 4.4: NQueen24

Figure 4.5: NQueen27

From these visualizations we can see that C, C++ and Rust performed with similar execution times, generally with less than 1 second every time. However, there was a clear difference in the execution times for AssemblyScript, which took between 3 and 5 seconds per execution.

4.2.2   Searching Algorithms

The source code for the Searching algorithms were kept syntactically the same as much as possible. An array of 100,000 elements was hard-coded into each of the algorithms. This caused the WebAssembly compilers generated huge WASM and WAT files for each language.

BinarySearch

For the BinarySearch algorithm, all languages had visually similar execution times, but it is still clear that there are differences between their performances. The statistical tests for C and C++ returned a value greater than 0.05, meaning that there was no significant difference between their execution times. However, Rust appears to have the fastest execution times over C and C++, while AssemblyScript has once again the slowest execution times. There were also significant differences in the inspection of the WAT files for the BinarySearch algorithm when converted into WebAssembly. The execution times for the BinarySearch algorithm can be view in figure 4.6.

Figure 4.6: BinarySearch

LinearSearch

The LinearSearch algorithm statistical tests showed that all languages had significant differences in their execution times, although visually, C and C++ appear to have similar performances. Once again Rust had the fastest execution times, generally running at around 4 seconds each time. C and C++ executed at about 5 to 6 seconds each time, while AssemblyScript had the slowest execution times, performing between 12 and 16 seconds for each iteration. The figure 4.7 shows the execution times for the LinearSearch algorithm in the selected languages.

Figure 4.7: LinearSearch

4.2.3 Sorting Algorithms

From the following visualizations for the sorting algorithms, it can be determined that C and C++ consistently had the fastest execution times over the other languages, while Rust and AssemblyScript had a mixed result, with sometimes one performing better than the other.

HeapSort

According to the statistical tests we ran on the HeapSort algorithm, C and C++ had a p-value of greater than 0.05, meaning that there was no significant difference between their execution times. This can be visually verified in figure 4.8. An interesting observation is that Rust appears to start slowly and then speed up during the experiment. This experiment was repeated many times and the same behavior from Rust was always observed. The results show that C and C++ had the fastest times, followed by Rust, and AssemblyScript had the slowest times.

Figure 4.8: HeapSort

MergeSort

The statistical tests performed on the MergeSort algorithm also returned a p-value greater than 0.05 for C and C++, meaning there was no significant difference between their execution times. This is visible on figure 4.9. Here we can see that once again Rust starts off slower than begins to speed up. Another interesting observation is that AssemblyScript actually runs faster than Rust for this algorithm, while C and C++ have once again the fastest execution times.

Figure 4.9: MergeSort

ShellSort

The p-values for the ShellSort algorithm for C and C++ again came back with values greater than 0.05, meaning that there was no significant difference between their execution times. This can be visually confirmed on figure 4.10. It can also be observed that Rust once again starts off slower, then speeds up towards the end of the execution. C and C++ have the fastest execution times, while Rust and AssemblyScript have visually similar execution times.

Figure 4.10: ShellSort

SelectionSort

The SelectionSort visualization shows that C and C++ have visually similar performances, however the statistical tests returned a p-value of less than 0.05. From this visualization we can see that C and C++ generally took around 4 seconds to execute, while Rust took around 8 seconds. The performance of AssemblyScript however was much slower, taking between 30 and 40 seconds for each iteration.

Figure 4.11: SelectionSort

BubbleSort

The BubbleSort algorithm revealed great differences between the performance of AssemblyScript and the other languages. While C and C++ performed slightly faster than Rust, with execution times of around 7 to 8 seconds, and Rust with execution times of around 9 to 10 seconds, AssemblyScript had taken between 70 and 80 seconds to perform each iteration.

Figure 4.12: BubbleSort

4.3   Discussion

WebAssembly is generally considered to be still in its infancy stages, and many features are still being developed. This imposes some restrictions on what compilers can actually compile into WebAssembly at this time. Overall, the performance of AssemblyScript was considerably slower than Rust, C and C++. The performance

times of C and C++ were mostly faster than all of the other languages. However, there were times that Rust performed faster than C/C++ and there was one occasion where AssemblyScript performed faster than Rust. Some of the statistical test results showed that there were significant difference between C and C++, however we believe that given a much higher number of iterations in each experiment, that the statistical tests would eventually show no significant difference between all of the algorithms implemented in C and C++.

Although not within the scope of this research, there was another interesting observation made regarding the compilers. The C, C++ and AssemblyScript modules usually compiled with reasonable speeds, normally ranging between 1 and 5 seconds for each algorithm. However, for the algorithms with the arrays of 100,000 elements, the Rust compiler would take around 30 minutes to compile.

The research work done in this paper highlights an important aspect of WebAssembly compilers, in that they do in fact suffer from significantly different performance times, even if the source code is syntactically identical in the various different programming languages. The idea behind WebAssembly is to allow multiple languages, such as C and Rust, to be used on the web. However, since the compilers will ultimately produce different WebAssembly files with varying performances, developers should give careful consideration into which language and compiler they use to generate WebAssembly modules.

Given more time to complete this research paper, we would have included more WebAssembly compilers in our experiments and evaluations. Also, the number of iterations for each algorithm would ideally be increased in order to produce a normal distribution of data and to use parametric tests to either accept of reject the hypothesis. Another aspect that could be improved is the choice of algorithms for each language. We believe that the Ostrich Benchmark Suite would be a suitable candidate. Ideally these benchmarks would be written in a variety of programming languages, such as Haskell and Prolog, covering a variety of programming paradigms such as procedural, object-oriented, function and logical. We believe the analysis of the same algorithms written across a large spectrum of programming languages and paradigms converted

into WebAssembly would provide interesting and useful insights.

One final observation that we will make in our evaluation is the comparison of the total time in milliseconds that each language took to execute each of the algorithms for every iteration. This resulted in C++ being the fastest overall, followed C and Rust, with AssemblyScript being the slowest overall. Also worth mentioning is that the Emscripten compiler is actually capable of achieving even higher optimization levels than the level -O2 that we used in our experiments, at the cost of longer compile times. However, even so, C/C++ compiled into WebAssembly with Emscripten provided the best overall performance against the other languages and compilers we evaluated. This reveals that currently, although WebAssembly promises speeds of near native performance otherwise only achieved with C/C++, that WebAssembly compiled from C/C++ still provide the best performances over other languages compiled into WebAssembly. This can be viewed in figure 4.13.

	1st	2nd	3rd	4th
	CPP	C	Rust	AssemblyScript
Fibonacci	813.9275436	832.3730369	747.5073359	1028.481608
PrimeNumber	59.14979537	58.05474328	112.3503863	547.5629386
NQueen24	36.50113993	43.55451636	45.26297373	376.1207794
NQueen27	50.41014082	58.15290738	63.60778665	517.3671882
BinarySearch	3.634928785	3.55388708	2.948940285	4.70987111
LinearSearch	653.9630445	673.3244232	432.8657695	1627.994193
HeapSort	2.20110994	2.192816007	3.611971927	14.57914214
MergeSort	1.44226305	1.411856896	9.854781747	6.714673676
ShellSort	1.071705701	1.047660574	2.829952329	3.421946958
SelectionSort	486.0118547	499.6437487	885.545482	3864.948572
BubbleSort	943.4891187	1010.069966	1158.012824	8793.64782
	3051.802645	3183.379562	3464.398204	16785.54873

Figure 4.13: Overall Performance in milliseconds

Chapter 5

Conclusion

5.1 Research Overview

WebAssembly is becoming more and more popular in research and web-development as it promises improved performance to compute-intensive operations for web-based applications. In this paper we critically evaluate the current literature that benchmarks WebAssembly performance compared with traditional implementations. In our opinion, there is no doubt that WebAssembly offers an exciting and efficient alternative to many of the limitations found in web-based and cloud-based applications of today.

However, we raise the concern that the performance of WebAssembly is greatly impacted by the choice of compiler used, meaning that certain programming languages will still have significant differences in their performance when compiled to WebAssembly. At the time of writing this paper, we did not find any other research paper that compares the code generated by WebAssembly compilers and evaluates their performance against each other.

We also highlight that many research papers have proposed WebAssembly alternatives to native applications, performing benchmarks to evaluate the performance of WebAssembly. However, these evaluations only ever used one compiler for the generation of WebAssembly. Based on the results of our experiments, we show that different WebAssembly compilers produce different results in terms of performance. There-

fore we believe that evaluating the performance of WebAssembly generated from only one compiler is not sufficient when performing evaluations between WebAssembly and native speeds.

5.2 Problem Definition

The objective of this research was to evaluate the performance of WebAssembly code generated by different compilers, highlighting a potential issue with the performance of the generated code. For example, if we write even a simple loop in C, C++, Rust and AssemblyScript, one might initially assume that the WebAssembly code produced by the compilers would be similar and would run without significant differences in performance. However, we have highlighted that this is not the case, demonstrating that each of the compilers produced significantly different results in terms of execution speed.

5.3 Design/Experimentation, Evaluation & Results

For our experiments, we used a selection of 11 different programs written in C, C++, Rust and AssemblyScript. These programs ranged from simple Fibonacci, Prime Number, Searching and Sorting functions to more demanding numerical computing functions such as the N-Queens problem. We compiled each program into WebAssembly using the respective compiler for each programming language.

We used NodeJS to run the WebAssembly compiled functions and measure their execution times. Each function was run a total of 120 times. Upon inspection of the data produced by these experiments, we determined that the data did not have a normal distribution. This in turn influenced the choice of statistical test that we had to use to determine whether or not there were significant differences in the performance of the compiled WebAssembly modules.

Our results showed that in fact there are significant differences in the performance of WebAssembly modules generated by the different compilers, not only in terms of

execution times, but upon visual inspection of the human-readable WAT format of the WebAssembly code, we can clearly see significant differences even for simple programs compiled to WebAssembly.

We also revealed that, although WebAssembly promises performance of near native speed otherwise only achieved with C/C++, currently WebAssembly compiled from C/C++ still offer the best performance over the other programming languages used in our research.

5.4   Contributions and impact

In this research we have identified a potential problem with WebAssembly compilers, in that they produce different byte code for even simple programs, and have significant differences in their performance, depending on which compiler was used to generate the WebAssembly module. We show that even though WebAssembly promises performances of near native speeds otherwise only obtainable with C/C++ code, that when C/C++ code is compiled into WebAssembly, it still out-performes all the other languages in our experiments. We believe that this provides valuable insights for developers as they choose which programming language to develop in and subsequently compile into WebAssembly.

We also provided a critical evaluation of the literature currently available for WebAssembly benchmarking and for WebAssembly solutions being proposed as alternatives to existing approaches in applications. We revealed a gap in the current literature and highlighted that for evaluating the performance of WebAssembly, using only one compiler is not sufficient, as different WebAssembly compilers may produce different results, which in turn may determine the success or failure of the WebAssembly evaluations currently being researched.

5.5   Future Work & recommendations

In chapter 1 of this paper we identified the scope and limitations of our research. These limitations create opportunities for future work in this area. We recommend that more research be done on the comparison of code generated by WebAssembly compilers to identify the differences between their produced code and determine the reasons why these differences exist. Another area of interest could be to compare the file size of the compiled WebAssembly modules from different compilers while compiling the same programming logic.

The list of programming languages that currently compile to WebAssembly can be found online[1]. Future research should be done to evaluate the performance of all available WebAssembly compilers. Furthermore, as more and more compilers emerge, programming languages with different programming paradigms will be able to compile to WebAssembly, for example Haskell, a functional programming language and ProLog, a logic programming language. We believe that analysing the WebAssembly byte code and performance of functional, logical, procedural and object-oriented programming languages will be of interest to researchers and application developers.

The Ostrich Benchmark Suite was published in (Khan et al., 2015). They identified important patterns for numerical computation and ran experiments to compare JavaScript performance against native C code for sequential and parallel operations. Therefore another possibility for future work would be to implement the numerical computing benchmarks identified in the Ostrich Benchmark Suite in multiple programming languages, including languages from different programming paradigms such as procedural, functional, logical and object-oriented, and evaluating the performance of these benchmarks in WebAssembly from multiple compilers.

Finally, since our experiments were only executed on a NodeJS environment, future work might also involve running such experiments on all web-browsers and IoT devices.

---

[1]https://github.com/appcypher/awesome-wasm-langs

Bibliography

Abhay, G., Abhishek, S., & Namita, G. (2019). A variant of Bucket Sort Shell Sort vs Insertion Sort. *2019 10th International Conference on Computing, Communication and Networking Technologies, ICCCNT 2019*, 6–10. `https://doi.org/10.1109/ICCCNT45670.2019.8944607`

Alhassan, A. (2019). Build and conquer: Solving N queens problem using iterative compression. *Proceedings of the International Conference on Computer, Control, Electrical, and Electronics Engineering 2019, ICCCEEE 2019*. `https://doi.org/10.1109/ICCCEEE46830.2019.9070976`

Ayub, M. A., Kalpoma, K. A., Proma, H. T., Kabir, S. M., & Chowdhury, R. I. H. (2018). Exhaustive study of essential constraint satisfaction problem techniques based on N-Queens problem. *20th International Conference of Computer and Information Technology, ICCIT 2017, 2018-Janua*, 1–6. `https://doi.org/10.1109/ICCITECHN.2017.8281850`

Cabrera Arteaga, J., Donde, S., Gu, J., Floros, O., Satabin, L., Baudry, B., & Monperrus, M. (2020). Superoptimization of WebAssembly bytecode. *ACM International Conference Proceeding Series*, 36–40. `https://doi.org/10.1145/3397537.3397567`

Elhakeem Abd Elnaby, A., & El-Baz, A. H. (2021). A new explicit algorithmic method for generating the prime numbers in order. *Egyptian Informatics Journal, 22*(1), 101–104. `https://doi.org/10.1016/j.eij.2020.05.002`

Guldal, S., Baugh, V., & Allehaibi, S. (2016). N-Queens solving algorithm by sets and backtracking. *Conference Proceedings - IEEE SOUTHEASTCON, 2016-July*. `https://doi.org/10.1109/SECON.2016.7506688`

Haas, A., Rossberg, A., Schuff, D. L., Titzer, B. L., Holman, M., Gohman, D., Wagner, L., Zakai, A., & Bastien, J. F. (2017). Bringing the web up to speed with WebAssembly. *ACM SIGPLAN Notices*, *52*(6), 185–200. `https://doi.org/10.1145/3062341.3062363`

Hall, A., & Ramachandran, U. (2019). An execution model for serverless functions at the edge. *IoTDI 2019 - Proceedings of the 2019 Internet of Things Design and Implementation*, 225–236. `https://doi.org/10.1145/3302505.3310084`

Herrera, D., Chen, H., Lavoie, E., & Hendren, L. (2018). WebAssembly and JavaScript Challenge: Numerical program performance using modern browser technologies and devices.

Jacob, A. E., Ashodariya, N., & Dhongade, A. (2018). Hybrid search algorithm: Combined linear and binary search algorithm. *2017 International Conference on Energy, Communication, Data Analytics and Soft Computing, ICECDS 2017*, 1543–1547. `https://doi.org/10.1109/ICECDS.2017.8389704`

Jangda, A., Powers, B., Berger, E. D., & Guha, A. (2019). Not so fast: Analyzing the performance of webassembly vs. native code. *2019 USENIX Annual Technical Conference (USENIX ATC 19)*, 107–120. `https://www.usenix.org/conference/atc19/presentation/jangda`

Jeong, H. J., Shin, C. H., Shin, K. Y., Lee, H. J., & Moon, S. M. (2019). Seamless Offloading of Web App Computations from Mobile Device to Edge Clouds via HTML5 Web Worker Migration. *SoCC 2019 - Proceedings of the ACM Symposium on Cloud Computing*, 38–49. `https://doi.org/10.1145/3357223.3362735`

Khan, F., Foley-Bourgon, V., Kathrotia, S., Lavoie, E., & Hendren, L. (2015). Using JavaScript and WebCL for numerical computations: A comparative study of native and web technologies. *ACM SIGPLAN Notices*, *50*(2), 91–102. `https://doi.org/10.1145/2661088.2661090`

Koren, I. (2021). A Standalone WebAssembly Development Environment for the Internet of Things, 353–360. `https://doi.org/10.1007/978-3-030-74296-6_27`

Letz, S., Orlarey, Y., & Fober, D. (2018). FAUST Domain Specific Audio DSP Language Compiled to WebAssembly. *The Web Conference 2018 - Companion of the World Wide Web Conference, WWW 2018*, 701–709. `https://doi.org/10.1145/3184558.3185970`

Long, J., Tai, H. Y., Hsieh, S. T., & Yuan, M. J. (2021). A Lightweight Design for Serverless Function as a Service. *IEEE Software*, *38*(1), 75–80. `https://doi.org/10.1109/MS.2020.3028991`

Mendki, P. (2020). Evaluating webassembly enabled serverless approach for edge computing. *Proceedings - 2020 IEEE Cloud Summit, Cloud Summit 2020*, 161–166. `https://doi.org/10.1109/IEEECloudSummit48914.2020.00031`

Murphy, S., Persaud, L., Martini, W., & Bosshard, B. (2020). On the use of web assembly in a serverless context. *Lecture Notes in Business Information Processing*, *396 LNBIP*, 141–145. `https://doi.org/10.1007/978-3-030-58858-8_15`

Musch, M., Wressnegger, C., Johns, M., & Rieck, K. (2019). New kid on the web: A study on the prevalence of webassembly in the wild. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, *11543 LNCS*, 23–42. `https://doi.org/10.1007/978-3-030-22038-9_2`

Protzenko, J., Beurdouche, B., Merigoux, D., & Bhargavan, K. (2019). Formally verified cryptographic web applications in webassembly. *Proceedings - IEEE Symposium on Security and Privacy*, *2019-May*, 1256–1274. `https://doi.org/10.1109/SP.2019.00064`

Reiser, M., & Bläser, L. (2017). Accelerate javascript applications by cross-compiling towebassembly. *VMIL 2017 - Proceedings of the 9th ACM SIGPLAN International Workshop on Virtual Machines and Intermediate Languages, co-located with SPLASH 2017*, 10–17. `https://doi.org/10.1145/3141871.3141873`

Sandhu, P., Herrera, D., & Hendren, L. (2018). Sparse matrices on the web - Characterizing the performance and optimal format selection of sparse matrix-vector multiplication in JavaScript and WebAssembly. *ACM International Conference Proceeding Series*. `https://doi.org/10.1145/3237009.3237020`

Sandhu, P., Verbrugge, C., & Hendren, L. (2020). A fully structure-driven performance analysis of sparse matrix-vector multiplication. *ICPE 2020 - Proceedings of the ACM/SPEC International Conference on Performance Engineering*, 108–119. https://doi.org/10.1145/3358960.3379131

Shillaker, S., & Pietzuch, P. (2020). FAASM: Lightweight isolation for efficient stateful serverless computing. *Proceedings of the 2020 USENIX Annual Technical Conference, ATC 2020*, 419–433.

Taheri, S. (2018). OpenCV . js : Computer Vision Processing for the Open Web Platform.

Tiwary, M., Mishra, P., Jain, S., & Puthal, D. (2020). Data Aware Web-Assembly Function Placement. *The Web Conference 2020 - Companion of the World Wide Web Conference, WWW 2020*, 4–5. https://doi.org/10.1145/3366424 .3382670

van Eekelen, M., Frumin, D., Geuvers, H., Gondelman, L., Krebbers, R., Schoolderman, M., Smetsers, S., Verbeek, F., Viguier, B., & Wiedijk, F. (2019). A benchmark for C program verification. *arXiv*.

Watt, C., Renner, J., Popescu, N., Cauligi, S., & Stefan, D. (2019). CT-wasm: type-driven secure cryptography for the web ecosystem. *Proceedings of the ACM on Programming Languages*, *3*(POPL), 1–29. https://doi.org/10.1145/ 3290390

Zheng, S., Wang, H., Wu, L., Huang, G., & Liu, X. (2020). VM matters: A comparison of WASM vms and evms in the performance of blockchain smart contracts. *CoRR*, *abs/2012.01032*. https://arxiv.org/abs/2012.01032

Appendix A

Source Code

The entire code and results for these experiments can be found at:

https://github.com/rayphelan/MastersProject2021

Listing A.1 contains the entire RStudio code used for performing statistical tests.

Listing A.1: Statistical Tests in RStudio

```
 1  library("dplyr")
 2  library("ggpubr")
 3
 4  ################################################################################
 5  # Load Languages
 6  AssemblyScript <- read.csv('C:/TuDublin/Masters Project/Git Clone/MastersProject2021/
 7  FinalResultsPerLanguage/CSV/AssemblyScript.csv')
 8
 9  C <- read.csv('C:/TuDublin/Masters Project/Git Clone/MastersProject2021/
10  FinalResultsPerLanguage/CSV/C.csv')
11
12  CPP <- read.csv('C:/TuDublin/Masters Project/Git Clone/MastersProject2021/
13  FinalResultsPerLanguage/CSV/CPP.csv')
14
15  Rust <- read.csv('C:/TuDublin/Masters Project/Git Clone/MastersProject2021/
16  FinalResultsPerLanguage/CSV/Rust.csv')
17
18
19  ################################################################################
20  # Density Plot - AssemblyScript
```

```
21   ggdensity ( AssemblyScript$BinarySearch )
22   ggdensity ( AssemblyScript$BubbleSort )
23   ggdensity ( AssemblyScript$Fibonacci )
24   ggdensity ( AssemblyScript$HeapSort )
25   ggdensity ( AssemblyScript$LinearSearch )
26   ggdensity ( AssemblyScript$MergeSort )
27   ggdensity ( AssemblyScript$NQueen24 )
28   ggdensity ( AssemblyScript$NQueen27 )
29   ggdensity ( AssemblyScript$PrimeNumber )
30   ggdensity ( AssemblyScript$SelectionSort )
31   ggdensity ( AssemblyScript$ShellSort )
32
33   # Normality Plot − AssemblyScript
34   ggqqplot ( AssemblyScript$BinarySearch )
35   ggqqplot ( AssemblyScript$BubbleSort )
36   ggqqplot ( AssemblyScript$Fibonacci )
37   ggqqplot ( AssemblyScript$HeapSort )
38   ggqqplot ( AssemblyScript$LinearSearch )
39   ggqqplot ( AssemblyScript$MergeSort )
40   ggqqplot ( AssemblyScript$NQueen24 )
41   ggqqplot ( AssemblyScript$NQueen27 )
42   ggqqplot ( AssemblyScript$PrimeNumber )
43   ggqqplot ( AssemblyScript$SelectionSort )
44   ggqqplot ( AssemblyScript$ShellSort )
45
46   # Histograms − AssemblyScript
47   hist ( AssemblyScript$BinarySearch )
48   hist ( AssemblyScript$BubbleSort )
49   hist ( AssemblyScript$Fibonacci )
50   hist ( AssemblyScript$HeapSort )
51   hist ( AssemblyScript$LinearSearch )
52   hist ( AssemblyScript$MergeSort )
53   hist ( AssemblyScript$NQueen24 )
54   hist ( AssemblyScript$NQueen27 )
55   hist ( AssemblyScript$PrimeNumber )
56   hist ( AssemblyScript$SelectionSort )
57   hist ( AssemblyScript$ShellSort )
58
59   # Shapiro−Wilk normality test − AssemblyScript
60   shapiro.test ( AssemblyScript$BinarySearch )    # W = 0.89333,  p−value = 9.081e−08
61   shapiro.test ( AssemblyScript$BubbleSort )      # W = 0.947,    p−value = 0.000131
62   shapiro.test ( AssemblyScript$Fibonacci )       # W = 0.9183,   p−value = 1.921e−06
63   shapiro.test ( AssemblyScript$HeapSort )        # W = 0.92146,  p−value = 2.93e−06
64   shapiro.test ( AssemblyScript$LinearSearch )    # W = 0.85618,  p−value = 1.97e−09
65   shapiro.test ( AssemblyScript$MergeSort )       # W = 0.91813,  p−value = 1.878e−06
```

```
66   shapiro.test(AssemblyScript$NQueen24)              # W = 0.95549,   p-value = 0.0005585
67   shapiro.test(AssemblyScript$NQueen27)              # W = 0.9585,    p-value = 0.0009603
68   shapiro.test(AssemblyScript$PrimeNumber)           # W = 0.91855,   p-value = 1.986e-06
69   shapiro.test(AssemblyScript$SelectionSort)         # W = 0.73501,   p-value = 1.995e-13
70   shapiro.test(AssemblyScript$ShellSort)             # W = 0.89884,   p-value = 1.713e-07
71
72   # independent 2-group Mann-Whitney U Tests
73
74   # BinarySearch
75   wilcox.test(AssemblyScript$BinarySearch,C$BinarySearch)     # W = 12582,   p-value < 2.2e-16
76   wilcox.test(AssemblyScript$BinarySearch,CPP$BinarySearch)   # W = 12408,   p-value < 2.2e-16
77   wilcox.test(AssemblyScript$BinarySearch,Rust$BinarySearch)  # W = 13720,   p-value < 2.2e-16
78
79   # BubbleSort
80   wilcox.test(AssemblyScript$BubbleSort,C$BubbleSort)         # W = 14400,   p-value < 2.2e-16
81   wilcox.test(AssemblyScript$BubbleSort,CPP$BubbleSort)       # W = 14400,   p-value < 2.2e-16
82   wilcox.test(AssemblyScript$BubbleSort,Rust$BubbleSort)      # W = 14400,   p-value < 2.2e-16
83
84   # Fibonacci
85   wilcox.test(AssemblyScript$Fibonacci,C$Fibonacci)          # W = 14400,   p-value < 2.2e-16
86   wilcox.test(AssemblyScript$Fibonacci,CPP$Fibonacci)        # W = 14400,   p-value < 2.2e-16
87   wilcox.test(AssemblyScript$Fibonacci,Rust$Fibonacci)       # W = 14400,   p-value < 2.2e-16
88
89   # HeapSort
90   wilcox.test(AssemblyScript$HeapSort,C$HeapSort)            # W = 14400,   p-value < 2.2e-16
91   wilcox.test(AssemblyScript$HeapSort,CPP$HeapSort)          # W = 14400,   p-value < 2.2e-16
92   wilcox.test(AssemblyScript$HeapSort,Rust$HeapSort)         # W = 14400,   p-value < 2.2e-16
93
94   # LinearSearch
95   wilcox.test(AssemblyScript$LinearSearch,C$LinearSearch)     # W = 14400,   p-value < 2.2e-16
96   wilcox.test(AssemblyScript$LinearSearch,CPP$LinearSearch)   # W = 14400,   p-value < 2.2e-16
97   wilcox.test(AssemblyScript$LinearSearch,Rust$LinearSearch)  # W = 14400,   p-value < 2.2e-16
98
99   # MergeSort
100  wilcox.test(AssemblyScript$MergeSort,C$MergeSort)          # W = 14400,   p-value < 2.2e-16
101  wilcox.test(AssemblyScript$MergeSort,CPP$MergeSort)        # W = 14400,   p-value < 2.2e-16
102  wilcox.test(AssemblyScript$MergeSort,Rust$MergeSort)       # W = 362,     p-value < 2.2e-16
103
104  # NQueen24
105  wilcox.test(AssemblyScript$NQueen24,C$NQueen24)           # W = 14400,   p-value < 2.2e-16
106  wilcox.test(AssemblyScript$NQueen24,CPP$NQueen24)         # W = 14400,   p-value < 2.2e-16
107  wilcox.test(AssemblyScript$NQueen24,Rust$NQueen24)        # W = 14400,   p-value < 2.2e-16
108
109  # NQueen27
110  wilcox.test(AssemblyScript$NQueen27,C$NQueen27)           # W = 14400,   p-value < 2.2e-16
```

```
111   wilcox.test(AssemblyScript$NQueen27,CPP$NQueen27)              # W = 14400,   p-value < 2.2e-16
112   wilcox.test(AssemblyScript$NQueen27,Rust$NQueen27)             # W = 14400,   p-value < 2.2e-16
113
114   # PrimeNumber
115   wilcox.test(AssemblyScript$PrimeNumber,C$PrimeNumber)          # W = 14400,   p-value < 2.2e-16
116   wilcox.test(AssemblyScript$PrimeNumber,CPP$PrimeNumber)        # W = 14400,   p-value < 2.2e-16
117   wilcox.test(AssemblyScript$PrimeNumber,Rust$PrimeNumber)       # W = 14400,   p-value < 2.2e-16
118
119   # SelectionSort
120   wilcox.test(AssemblyScript$SelectionSort,C$SelectionSort)      # W = 14400, p-value < 2.2e-16
121   wilcox.test(AssemblyScript$SelectionSort,CPP$SelectionSort)    # W = 14400, p-value < 2.2e-16
122   wilcox.test(AssemblyScript$SelectionSort,Rust$SelectionSort)   # W = 14400, p-value < 2.2e-16
123
124   # ShellSort
125   wilcox.test(AssemblyScript$ShellSort,C$ShellSort)              # W = 14325,   p-value < 2.2e-16
126   wilcox.test(AssemblyScript$ShellSort,CPP$ShellSort)            # W = 14400,   p-value < 2.2e-16
127   wilcox.test(AssemblyScript$ShellSort,Rust$ShellSort)           # W = 9598,    p-value = 8.265e-06
128
129
130   ##########################################################################################
131   # Density Plot - C
132   ggdensity(C$BinarySearch)
133   ggdensity(C$BubbleSort)
134   ggdensity(C$Fibonacci)
135   ggdensity(C$HeapSort)
136   ggdensity(C$LinearSearch)
137   ggdensity(C$MergeSort)
138   ggdensity(C$NQueen24)
139   ggdensity(C$NQueen27)
140   ggdensity(C$PrimeNumber)
141   ggdensity(C$SelectionSort)
142   ggdensity(C$ShellSort)
143
144   # Normality Plot 1 - C
145   ggqqplot(C$BinarySearch)
146   ggqqplot(C$BubbleSort)
147   ggqqplot(C$Fibonacci)
148   ggqqplot(C$HeapSort)
149   ggqqplot(C$LinearSearch)
150   ggqqplot(C$MergeSort)
151   ggqqplot(C$NQueen24)
152   ggqqplot(C$NQueen27)
153   ggqqplot(C$PrimeNumber)
154   ggqqplot(C$SelectionSort)
155   ggqqplot(C$ShellSort)
```

```
156
157   # Histograms − C
158   hist(C$BinarySearch)
159   hist(C$BubbleSort)
160   hist(C$Fibonacci)
161   hist(C$HeapSort)
162   hist(C$LinearSearch)
163   hist(C$MergeSort)
164   hist(C$NQueen24)
165   hist(C$NQueen27)
166   hist(C$PrimeNumber)
167   hist(C$SelectionSort)
168   hist(C$ShellSort)
169
170   # Shapiro−Wilk normality test − C
171   shapiro.test(C$BinarySearch)      # W = 0.85759,  p−value = 2.251e−09
172   shapiro.test(C$BubbleSort)        # W = 0.90343,  p−value = 2.95e−07
173   shapiro.test(C$Fibonacci)         # W = 0.87602,  p−value = 1.394e−08
174   shapiro.test(C$HeapSort)          # W = 0.92663,  p−value = 5.966e−06
175   shapiro.test(C$LinearSearch)      # W = 0.92993,  p−value = 9.529e−06
176   shapiro.test(C$MergeSort)         # W = 0.92818,  p−value = 7.426e−06
177   shapiro.test(C$NQueen24)          # W = 0.95149,  p−value = 0.0002786
178   shapiro.test(C$NQueen27)          # W = 0.73445,  p−value = 1.927e−13
179   shapiro.test(C$PrimeNumber)       # W = 0.87064,  p−value = 8.051e−09
180   shapiro.test(C$SelectionSort)     # W = 0.91986,  p−value = 2.365e−06
181   shapiro.test(C$ShellSort)         # W = 0.79129,  p−value = 9.023e−12
182
183
184   # independent 2−group Mann−Whitney U Tests
185
186   # BinarySearch
187   wilcox.test(C$BinarySearch, AssemblyScript$BinarySearch)      # W = 1818,   p−value < 2.2e−16
188   wilcox.test(C$BinarySearch,CPP$BinarySearch)                  # W = 6410,   p−value = 0.1421
189   wilcox.test(C$BinarySearch, Rust$BinarySearch)               # W = 10409,  p−value = 2.427e−09
190
191   # BubbleSort
192   wilcox.test(C$BubbleSort, AssemblyScript$BubbleSort)          # W = 0,      p−value < 2.2e−16
193   wilcox.test(C$BubbleSort,CPP$BubbleSort)                      # W = 11118,  p−value = 3.225e−13
194   wilcox.test(C$BubbleSort, Rust$BubbleSort)                    # W = 626,    p−value < 2.2e−16
195
196   # Fibonacci
197   wilcox.test(C$Fibonacci, AssemblyScript$Fibonacci)            # W = 0,      p−value < 2.2e−16
198   wilcox.test(C$Fibonacci,CPP$Fibonacci)                        # W = 8162,   p−value = 0.07379
199   wilcox.test(C$Fibonacci, Rust$Fibonacci)                      # W = 12296,  p−value < 2.2e−16
200
```

```
201   # HeapSort
202   wilcox.test(C$HeapSort, AssemblyScript$HeapSort)           # W = 0,        p-value < 2.2e-16
203   wilcox.test(C$HeapSort, CPP$HeapSort)                      # W = 7423,     p-value = 0.6791
204   wilcox.test(C$HeapSort, Rust$HeapSort)                     # W = 3206,     p-value = 1.119e-13
205
206   # LinearSearch
207   wilcox.test(C$LinearSearch, AssemblyScript$LinearSearch)   # W = 0,        p-value < 2.2e-16
208   wilcox.test(C$LinearSearch, CPP$LinearSearch)              # W = 8953,     p-value = 0.001119
209   wilcox.test(C$LinearSearch, Rust$LinearSearch)             # W = 14280,    p-value < 2.2e-16
210
211   # MergeSort
212   wilcox.test(C$MergeSort, AssemblyScript$MergeSort)         # W = 0,        p-value < 2.2e-16
213   wilcox.test(C$MergeSort, CPP$MergeSort)                    # W = 6934,     p-value = 0.6215
214   wilcox.test(C$MergeSort, Rust$MergeSort)                   # W = 0,        p-value < 2.2e-16
215
216   # NQueen24
217   wilcox.test(C$NQueen24, AssemblyScript$NQueen24)           # W = 0,        p-value < 2.2e-16
218   wilcox.test(C$NQueen24, CPP$NQueen24)                      # W = 14090,    p-value < 2.2e-16
219   wilcox.test(C$NQueen24, Rust$NQueen24)                     # W = 5819,     p-value = 0.01026
220
221   # NQueen27
222   wilcox.test(C$NQueen27, AssemblyScript$NQueen27)           # W = 0,        p-value < 2.2e-16
223   wilcox.test(C$NQueen27, CPP$NQueen27)                      # W = 12789,    p-value < 2.2e-16
224   wilcox.test(C$NQueen27, Rust$NQueen27)                     # W = 2604,     p-value < 2.2e-16
225
226   # PrimeNumber
227   wilcox.test(C$PrimeNumber, AssemblyScript$PrimeNumber)     # W = 0,        p-value < 2.2e-16
228   wilcox.test(C$PrimeNumber, CPP$PrimeNumber)                # W = 4661,     p-value = 2.354e-06
229   wilcox.test(C$PrimeNumber, Rust$PrimeNumber)               # W = 0,        p-value < 2.2e-16
230
231   # SelectionSort
232   wilcox.test(C$SelectionSort, AssemblyScript$SelectionSort) # W = 0,        p-value < 2.2e-16
233   wilcox.test(C$SelectionSort, CPP$SelectionSort)            # W = 8881,     p-value = 0.001779
234   wilcox.test(C$SelectionSort, Rust$SelectionSort)           # W = 0,        p-value < 2.2e-16
235
236   # ShellSort
237   wilcox.test(C$ShellSort, AssemblyScript$ShellSort)         # W = 75,       p-value < 2.2e-16
238   wilcox.test(C$ShellSort, CPP$ShellSort)                    # W = 6754,     p-value = 0.4074
239   wilcox.test(C$ShellSort, Rust$ShellSort)                   # W = 529,      p-value < 2.2e-16
240
241
242   ################################################################################
243   # Density Plot - CPP
244   ggdensity(CPP$BinarySearch)
245   ggdensity(CPP$BubbleSort)
```

```
246   ggdensity (CPP$Fibonacci)
247   ggdensity (CPP$HeapSort)
248   ggdensity (CPP$LinearSearch)
249   ggdensity (CPP$MergeSort)
250   ggdensity (CPP$NQueen24)
251   ggdensity (CPP$NQueen27)
252   ggdensity (CPP$PrimeNumber)
253   ggdensity (CPP$SelectionSort)
254   ggdensity (CPP$ShellSort)
255
256   # Normality  Plot  1 − CPP
257   ggqqplot (CPP$BinarySearch)
258   ggqqplot (CPP$BubbleSort)
259   ggqqplot (CPP$Fibonacci)
260   ggqqplot (CPP$HeapSort)
261   ggqqplot (CPP$LinearSearch)
262   ggqqplot (CPP$MergeSort)
263   ggqqplot (CPP$NQueen24)
264   ggqqplot (CPP$NQueen27)
265   ggqqplot (CPP$PrimeNumber)
266   ggqqplot (CPP$SelectionSort)
267   ggqqplot (CPP$ShellSort)
268
269   # Histograms − CPP
270   hist (CPP$BinarySearch)
271   hist (CPP$BubbleSort)
272   hist (CPP$Fibonacci)
273   hist (CPP$HeapSort)
274   hist (CPP$LinearSearch)
275   hist (CPP$MergeSort)
276   hist (CPP$NQueen24)
277   hist (CPP$NQueen27)
278   hist (CPP$PrimeNumber)
279   hist (CPP$SelectionSort)
280   hist (CPP$ShellSort)
281
282   # Shapiro−Wilk  normality  test − CPP
283   shapiro.test (CPP$BinarySearch)      # W = 0.88697,   p−value = 4.471e−08
284   shapiro.test (CPP$BubbleSort)        # W = 0.86648,   p−value = 5.32e−09
285   shapiro.test (CPP$Fibonacci)         # W = 0.87035,   p−value = 7.82e−09
286   shapiro.test (CPP$HeapSort)          # W = 0.92485,   p−value = 4.656e−06
287   shapiro.test (CPP$LinearSearch)      # W = 0.94416,   p−value = 8.256e−05
288   shapiro.test (CPP$MergeSort)         # W = 0.85896,   p−value = 2.564e−09
289   shapiro.test (CPP$NQueen24)          # W = 0.56648,   p−value < 2.2e−16
290   shapiro.test (CPP$NQueen27)          # W = 0.80628,   p−value = 2.8e−11
```

```
291   shapiro.test(CPP$PrimeNumber)        # W = 0.76245,   p-value = 1.183e-12
292   shapiro.test(CPP$SelectionSort)      # W = 0.91616,   p-value = 1.451e-06
293   shapiro.test(CPP$ShellSort)          # W = 0.91794,   p-value = 1.832e-06
294
295   # independent 2-group Mann-Whitney U Tests
296
297   # BinarySearch
298   wilcox.test(CPP$BinarySearch, AssemblyScript$BinarySearch)    # W = 1992,    p-value < 2.2e-16
299   wilcox.test(CPP$BinarySearch, C$BinarySearch)                 # W = 7990,    p-value = 0.1421
300   wilcox.test(CPP$BinarySearch, Rust$BinarySearch)              # W = 10886,   p-value = 7.219e-12
301
302   # BubbleSort
303   wilcox.test(CPP$BubbleSort, AssemblyScript$BubbleSort)        # W = 0,       p-value < 2.2e-16
304   wilcox.test(CPP$BubbleSort, C$BubbleSort)                     # W = 3282,    p-value = 3.225e-13
305   wilcox.test(CPP$BubbleSort, Rust$BubbleSort)                  # W = 73,      p-value < 2.2e-16
306
307   # Fibonacci
308   wilcox.test(CPP$Fibonacci, AssemblyScript$Fibonacci)         # W = 0,       p-value < 2.2e-16
309   wilcox.test(CPP$Fibonacci, C$Fibonacci)                      # W = 6238,    p-value = 0.07379
310   wilcox.test(CPP$Fibonacci, Rust$Fibonacci)                  # W = 11568,   p-value = 4.606e-16
311
312   # HeapSort
313   wilcox.test(CPP$HeapSort, AssemblyScript$HeapSort)           # W = 0,       p-value < 2.2e-16
314   wilcox.test(CPP$HeapSort, C$HeapSort)                        # W = 6977,    p-value = 0.6791
315   wilcox.test(CPP$HeapSort, Rust$HeapSort)                     # W = 3111,    p-value = 2.901e-14
316
317   # LinearSearch
318   wilcox.test(CPP$LinearSearch, AssemblyScript$LinearSearch)   # W = 0,       p-value < 2.2e-16
319   wilcox.test(CPP$LinearSearch, C$LinearSearch)                # W = 5447,    p-value = 0.001119
320   wilcox.test(CPP$LinearSearch, Rust$LinearSearch)             # W = 14280,   p-value < 2.2e-16
321
322   # MergeSort
323   wilcox.test(CPP$MergeSort, AssemblyScript$MergeSort)         # W = 0,       p-value < 2.2e-16
324   wilcox.test(CPP$MergeSort, C$MergeSort)                      # W = 7466,    p-value = 0.6215
325   wilcox.test(CPP$MergeSort, Rust$MergeSort)                   # W = 0,       p-value < 2.2e-16
326
327   # NQueen24
328   wilcox.test(CPP$NQueen24, AssemblyScript$NQueen24)           # W = 0,       p-value < 2.2e-16
329   wilcox.test(CPP$NQueen24, C$NQueen24)                        # W = 310,     p-value < 2.2e-16
330   wilcox.test(CPP$NQueen24, Rust$NQueen24)                     # W = 781,     p-value < 2.2e-16
331
332   # NQueen27
333   wilcox.test(CPP$NQueen27, AssemblyScript$NQueen27)           # W = 0,       p-value < 2.2e-16
334   wilcox.test(CPP$NQueen27, C$NQueen27)                        # W = 1611,    p-value < 2.2e-16
335   wilcox.test(CPP$NQueen27, Rust$NQueen27)                     # W = 872,     p-value < 2.2e-16
```

```
336
337   # PrimeNumber
338   wilcox.test(CPP$PrimeNumber, AssemblyScript$PrimeNumber)      # W = 0,     p-value < 2.2e-16
339   wilcox.test(CPP$PrimeNumber, C$PrimeNumber)                   # W = 9739,  p-value = 2.354e-06
340   wilcox.test(CPP$PrimeNumber, Rust$PrimeNumber)                # W = 0,     p-value < 2.2e-16
341
342   # SelectionSort
343   wilcox.test(CPP$SelectionSort, AssemblyScript$SelectionSort)  # W = 0,     p-value < 2.2e-16
344   wilcox.test(CPP$SelectionSort, C$SelectionSort)               # W = 5519,  p-value = 0.001779
345   wilcox.test(CPP$SelectionSort, Rust$SelectionSort)            # W = 0,     p-value < 2.2e-16
346
347   # ShellSort
348   wilcox.test(CPP$ShellSort, AssemblyScript$ShellSort)          # W = 0,     p-value < 2.2e-16
349   wilcox.test(CPP$ShellSort, C$ShellSort)                       # W = 7646,  p-value = 0.4074
350   wilcox.test(CPP$ShellSort, Rust$ShellSort)                    # W = 565,   p-value < 2.2e-16
351
352
353   ##################################################################################
354   # Density Plot - Rust
355   ggdensity(Rust$BinarySearch)
356   ggdensity(Rust$BubbleSort)
357   ggdensity(Rust$Fibonacci)
358   ggdensity(Rust$HeapSort)
359   ggdensity(Rust$LinearSearch)
360   ggdensity(Rust$MergeSort)
361   ggdensity(Rust$NQueen24)
362   ggdensity(Rust$NQueen27)
363   ggdensity(Rust$PrimeNumber)
364   ggdensity(Rust$SelectionSort)
365   ggdensity(Rust$ShellSort)
366
367   # Normality Plot 1 - Rust
368   ggqqplot(Rust$BinarySearch)
369   ggqqplot(Rust$BubbleSort)
370   ggqqplot(Rust$Fibonacci)
371   ggqqplot(Rust$HeapSort)
372   ggqqplot(Rust$LinearSearch)
373   ggqqplot(Rust$MergeSort)
374   ggqqplot(Rust$NQueen24)
375   ggqqplot(Rust$NQueen27)
376   ggqqplot(Rust$PrimeNumber)
377   ggqqplot(Rust$SelectionSort)
378   ggqqplot(Rust$ShellSort)
379
380   # Histograms - Rust
```

71

```
381    hist ( Rust$BinarySearch )
382    hist ( Rust$BubbleSort )
383    hist ( Rust$Fibonacci )
384    hist ( Rust$HeapSort )
385    hist ( Rust$LinearSearch )
386    hist ( Rust$MergeSort )
387    hist ( Rust$NQueen24 )
388    hist ( Rust$NQueen27 )
389    hist ( Rust$PrimeNumber )
390    hist ( Rust$SelectionSort )
391    hist ( Rust$ShellSort )
392
393    # Shapiro−Wilk normality test − Rust
394    shapiro.test(Rust$BinarySearch)        # W = 0.9888,    p−value = 0.4335
395    shapiro.test(Rust$BubbleSort)          # W = 0.91152,   p−value = 8.002e−07
396    shapiro.test(Rust$Fibonacci)           # W = 0.90205,   p−value = 2.501e−07
397    shapiro.test(Rust$HeapSort)            # W = 0.90527,   p−value = 3.685e−07
398    shapiro.test(Rust$LinearSearch)        # W = 0.45708,   p−value < 2.2e−16
399    shapiro.test(Rust$MergeSort)           # W = 0.56386,   p−value < 2.2e−16
400    shapiro.test(Rust$NQueen24)            # W = 0.97237,   p−value = 0.01424
401    shapiro.test(Rust$NQueen27)            # W = 0.79029,   p−value = 8.382e−12
402    shapiro.test(Rust$PrimeNumber)         # W = 0.88119,   p−value = 2.399e−08
403    shapiro.test(Rust$SelectionSort)       # W = 0.88721,   p−value = 4.591e−08
404    shapiro.test(Rust$ShellSort)           # W = 0.97768,   p−value = 0.04334
405
406
407    # independent 2−group Mann−Whitney U Tests
408
409    # BinarySearch
410    wilcox.test(Rust$BinarySearch, AssemblyScript$BinarySearch)   # W = 680,    p−value < 2.2e−16
411    wilcox.test(Rust$BinarySearch,C$BinarySearch)                 # W = 3991,   p−value = 2.427e−09
412    wilcox.test(Rust$BinarySearch,CPP$BinarySearch)               # W = 3514,   p−value = 7.219e−12
413
414    # BubbleSort
415    wilcox.test(Rust$BubbleSort, AssemblyScript$BubbleSort)       # W = 0,      p−value < 2.2e−16
416    wilcox.test(Rust$BubbleSort,C$BubbleSort)                     # W = 13774,  p−value < 2.2e−16
417    wilcox.test(Rust$BubbleSort,CPP$BubbleSort)                   # W = 14327,  p−value < 2.2e−16
418
419    # Fibonacci
420    wilcox.test(Rust$Fibonacci, AssemblyScript$Fibonacci)         # W = 0,      p−value < 2.2e−16
421    wilcox.test(Rust$Fibonacci,C$Fibonacci)                       # W = 2104,   p−value < 2.2e−16
422    wilcox.test(Rust$Fibonacci,CPP$Fibonacci)                     # W = 2832,   p−value = 4.606e−16
423
424    # HeapSort
425    wilcox.test(Rust$HeapSort, AssemblyScript$HeapSort)           # W = 0,      p−value < 2.2e−16
```

```
426   wilcox.test(Rust$HeapSort,C$HeapSort)                              # W = 11194,  p−value = 1.119e−13
427   wilcox.test(Rust$HeapSort,CPP$HeapSort)                            # W = 11289,  p−value = 2.901e−14
428
429   # LinearSearch
430   wilcox.test(Rust$LinearSearch,AssemblyScript$LinearSearch)        # W = 0,      p−value < 2.2e−16
431   wilcox.test(Rust$LinearSearch,C$LinearSearch)                     # W = 120,    p−value < 2.2e−16
432   wilcox.test(Rust$LinearSearch,CPP$LinearSearch)                   # W = 120,    p−value < 2.2e−16
433
434   # MergeSort
435   wilcox.test(Rust$MergeSort,AssemblyScript$MergeSort)              # W = 14038,  p−value < 2.2e−16
436   wilcox.test(Rust$MergeSort,C$MergeSort)                           # W = 14400,  p−value < 2.2e−16
437   wilcox.test(Rust$MergeSort,CPP$MergeSort)                         # W = 14400,  p−value < 2.2e−16
438
439   # NQueen24
440   wilcox.test(Rust$NQueen24,AssemblyScript$NQueen24)                # W = 0,      p−value < 2.2e−16
441   wilcox.test(Rust$NQueen24,C$NQueen24)                             # W = 8581,   p−value = 0.01026
442   wilcox.test(Rust$NQueen24,CPP$NQueen24)                           # W = 13619,  p−value < 2.2e−16
443
444   # NQueen27
445   wilcox.test(Rust$NQueen27,AssemblyScript$NQueen27)                # W = 0,      p−value < 2.2e−16
446   wilcox.test(Rust$NQueen27,C$NQueen27)                             # W = 11796,  p−value < 2.2e−16
447   wilcox.test(Rust$NQueen27,CPP$NQueen27)                           # W = 13528,  p−value < 2.2e−16
448
449   # PrimeNumber
450   wilcox.test(Rust$PrimeNumber,AssemblyScript$PrimeNumber)          # W = 0,      p−value < 2.2e−16
451   wilcox.test(Rust$PrimeNumber,C$PrimeNumber)                       # W = 14400,  p−value < 2.2e−16
452   wilcox.test(Rust$PrimeNumber,CPP$PrimeNumber)                     # W = 14400,  p−value < 2.2e−16
453
454   # SelectionSort
455   wilcox.test(Rust$SelectionSort,AssemblyScript$SelectionSort)      # W = 0,      p−value < 2.2e−16
456   wilcox.test(Rust$SelectionSort,C$SelectionSort)                   # W = 14400,  p−value < 2.2e−16
457   wilcox.test(Rust$SelectionSort,CPP$SelectionSort)                 # W = 14400,  p−value < 2.2e−16
458
459   # ShellSort
460   wilcox.test(Rust$ShellSort,AssemblyScript$ShellSort)              # W = 4802,   p−value = 8.265e−06
461   wilcox.test(Rust$ShellSort,C$ShellSort)                           # W = 13871,  p−value < 2.2e−16
462   wilcox.test(Rust$ShellSort,CPP$ShellSort)                         # W = 13835,  p−value < 2.2e−16
463
464
465   ##################################################################################
466   ##################################################################################
467   ##################################################################################
```

Appendix B

Distribution of Data

The following figures represent the distribtion of the data obtained through statistical tests.

Figure B.1 shows the distribution of data for the Fibonacci sequence algorithm.


Figure B.1: Fibonacci

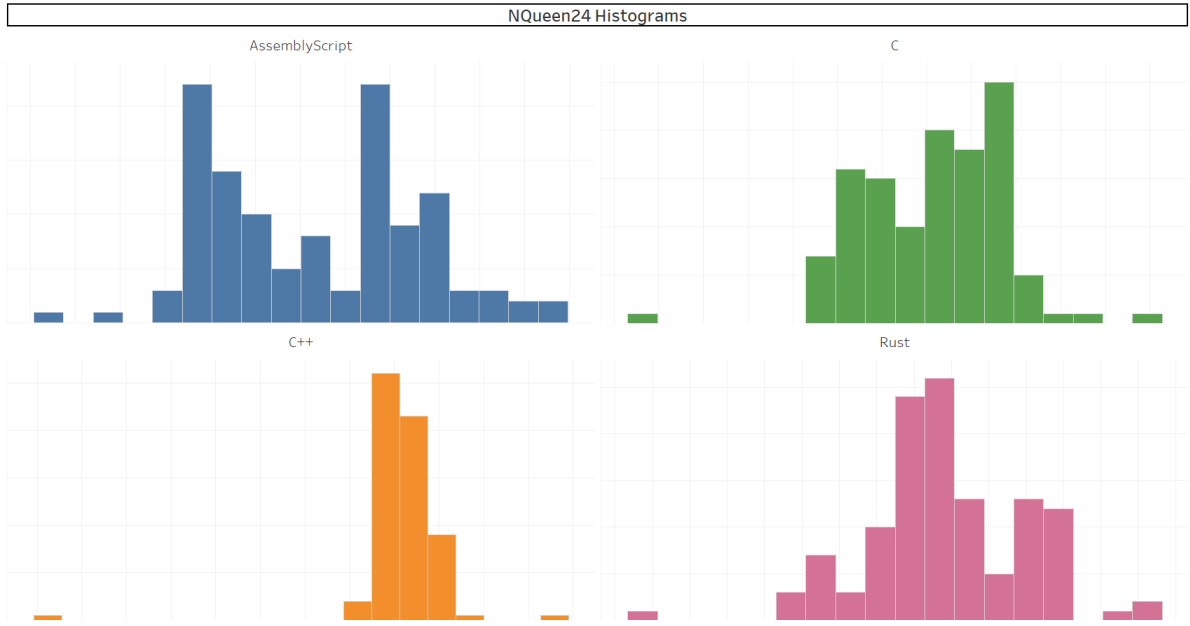Figure B.2 shows the distribution of data for the NQueen24 algorithm.

Figure B.2: NQueen24

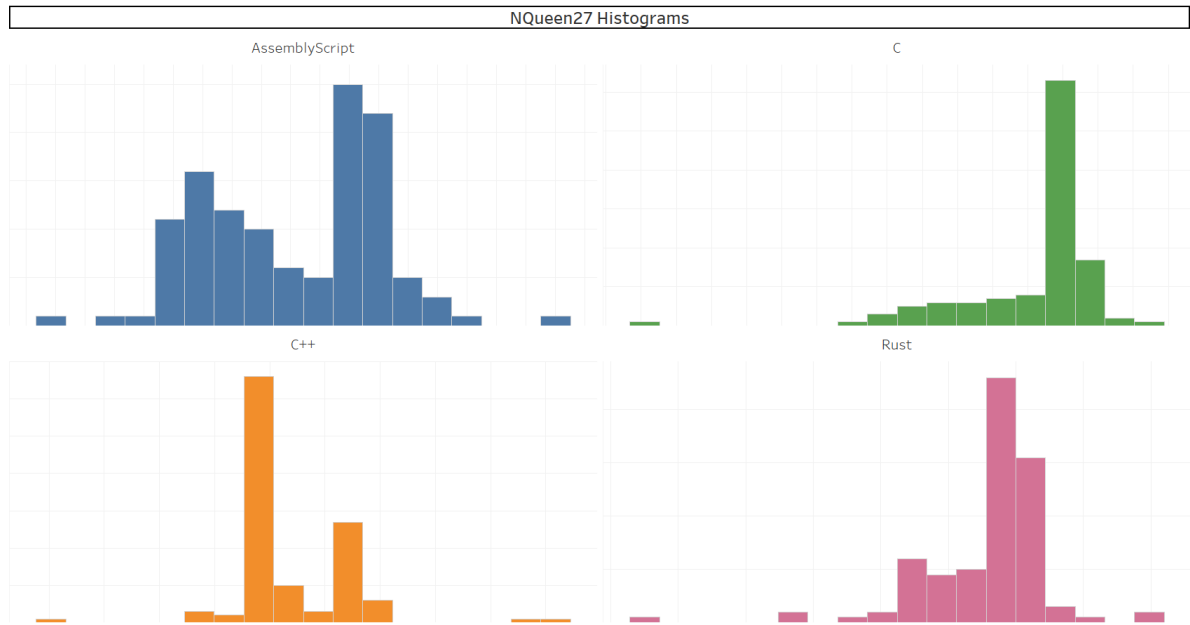Figure B.3 shows the distribution of data for the NQueen27 algorithm.

Figure B.3: NQueen27

Figure B.4 shows the distribution of data for the PrimeNumber algorithm.

Figure B.4: PrimeNumber

Figure B.5 shows the distribution of data for the BinarySearch algorithm.

Figure B.5: BinarySearch

Figure B.6 shows the distribution of data for the LinearSearch algorithm.

Figure B.6: LinearSearch

Figure B.7 shows the distribution of data for the BubbleSort algorithm.
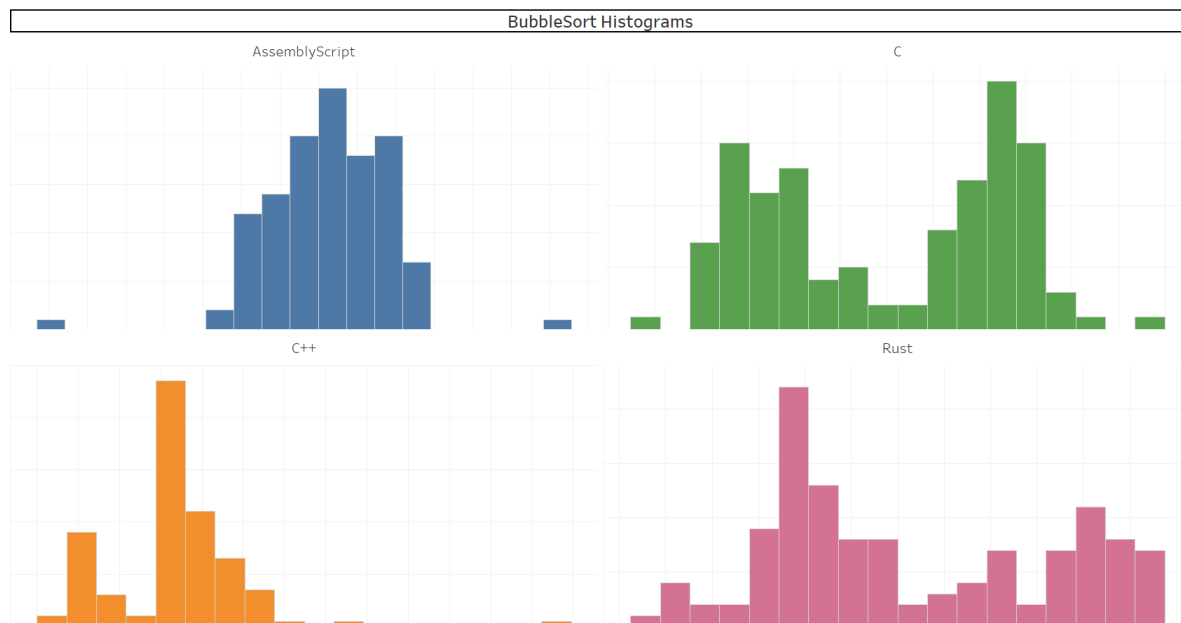
Figure B.7: BubbleSort

Figure B.8 shows the distribution of data for the HeapSort algorithm.


Figure B.8: HeapSort

Figure B.9 shows the distribution of data for the MergeSort algorithm.

Figure B.9: MergeSort

Figure B.10 shows the distribution of data for the ShellSort algorithm.


Figure B.10: ShellSort

Figure B.11 shows the distribution of data for the SelectionSort algorithm.


Figure B.11: SelectionSort

Figure B.12 represents the distribution of data for all the algorithms used in our experiments compiled from AssemblyScript.
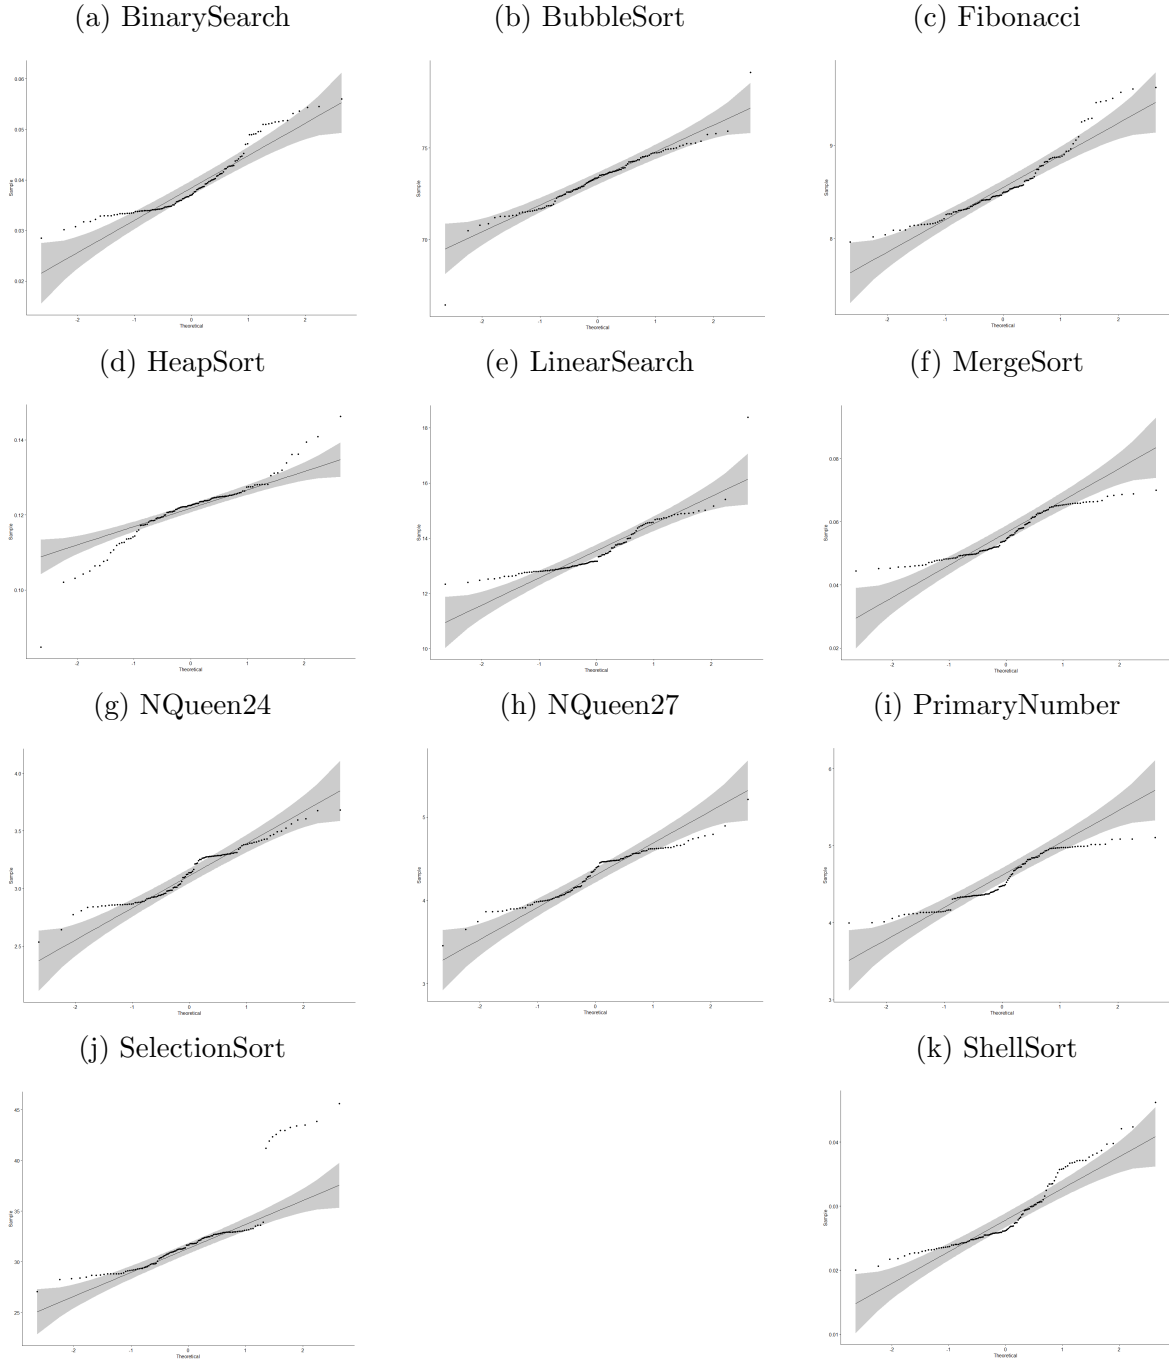
Figure B.12: AssemblyScript Distribtion

(a) BinarySearch (b) BubbleSort (c) Fibonacci

(d) HeapSort (e) LinearSearch (f) MergeSort

(g) NQueen24 (h) NQueen27 (i) PrimaryNumber

(j) SelectionSort (k) ShellSort

Figure 3.5 represents the distribution of data for all the algorithms used in our experiments compiled from C.

Figure B.13: C Distribtion

(a) BinarySearch

(b) BubbleSort

(c) Fibonacci

(d) HeapSort

(e) LinearSearch

(f) MergeSort

(g) NQueen24

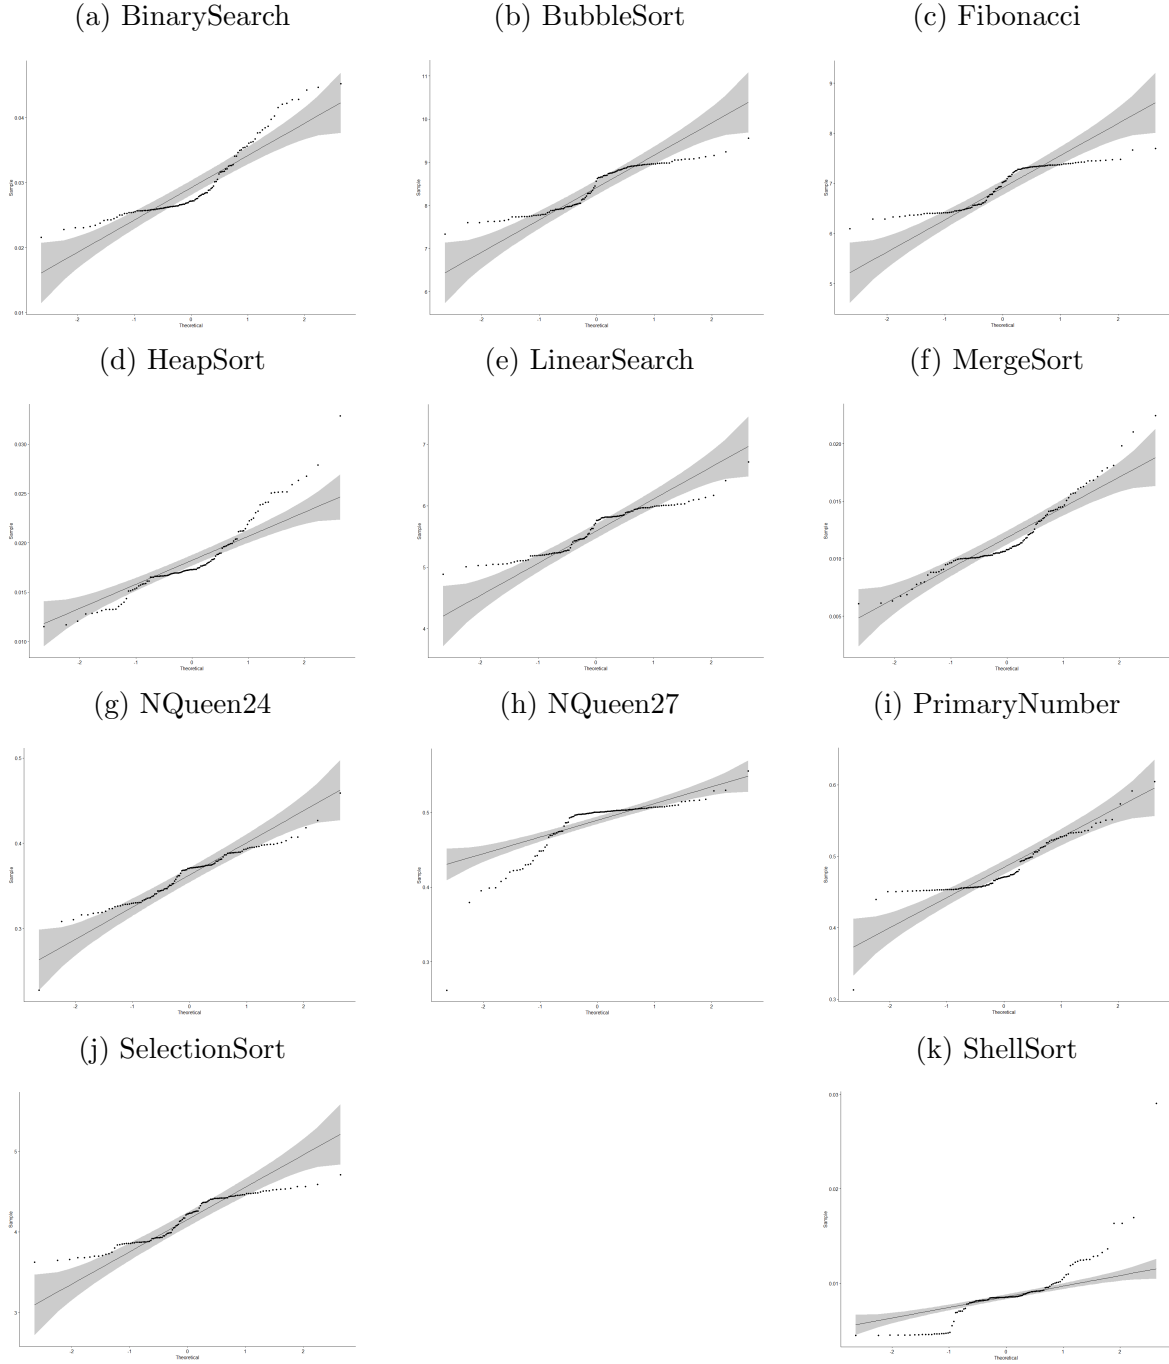(h) NQueen27

(i) PrimaryNumber

(j) SelectionSort

(k) ShellSort

Figure B.14 represents the distribution of data for all the algorithms used in our experiments compiled from C++.

Figure B.14: C++ Distribtion

(a) BinarySearch

(b) BubbleSort

(c) Fibonacci

(d) HeapSort

(e) LinearSearch

(f) MergeSort

(g) NQueen24

(h) NQueen27

(i) PrimaryNumber

(j) SelectionSort

(k) ShellSort

Figure B.15 represents the distribution of data for all the algorithms used in our experiments compiled from Rust.

Figure B.15: Rust Distribtion

(a) BinarySearch

(b) BubbleSort

(c) Fibonacci

(d) HeapSort

(e) LinearSearch

(f) MergeSort

(g) NQueen24

(h) NQueen27

(i) PrimaryNumber

(j) SelectionSort

(k) ShellSort